Getting Started with Amazon Redshift

Avinash Nidumbur
Sr Business Development Manager
Amazon Redshift
avinid@amazon.com

Agenda

- Introduction
- Benefits
- Use cases
- Getting started
- Q&A

Legacy architectural models lead to Dark Data

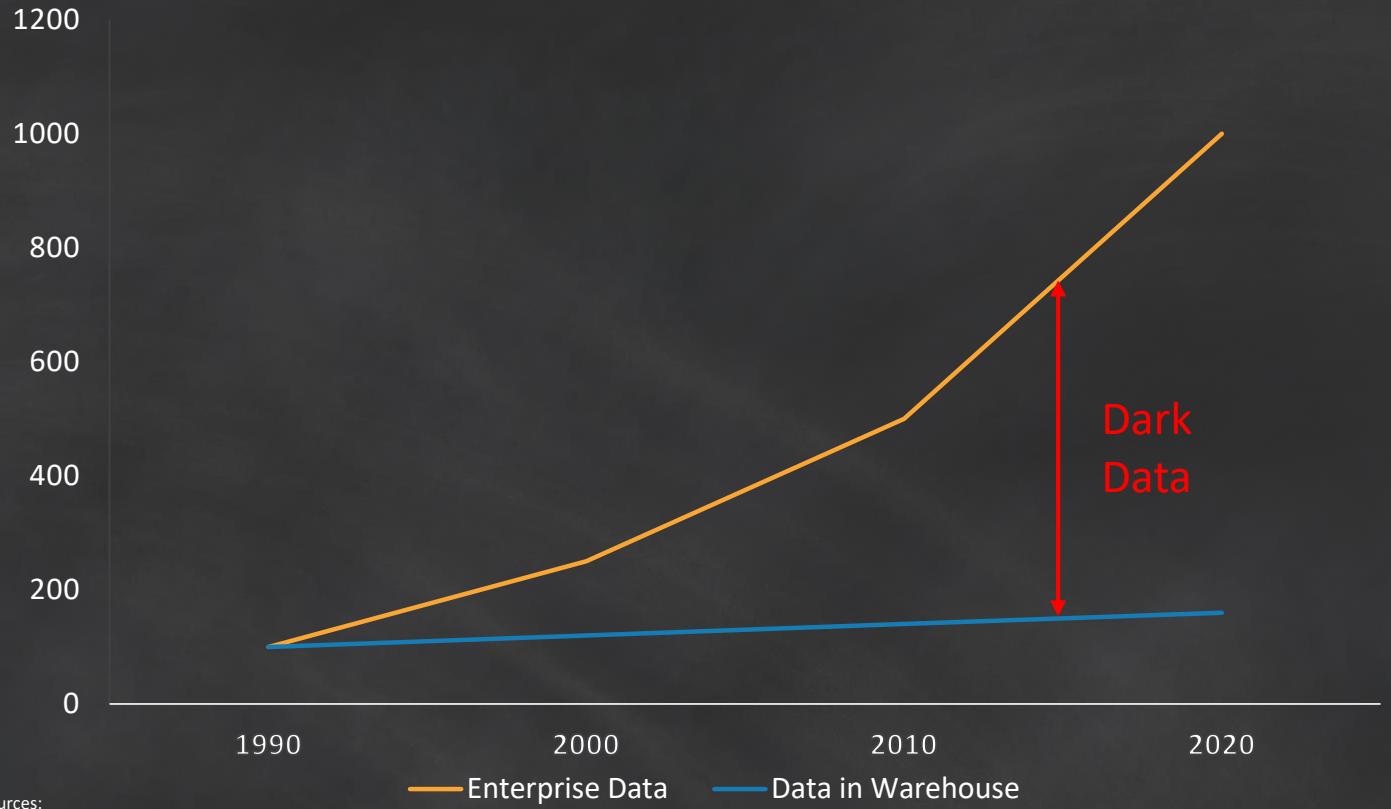

Very Expensive

Inflexible licensing

Proprietary

Lock-In

Sources:
Gartner: User Survey Analysis: Key Trends Shaping the Future of Data Center Infrastructure Through 2011
IDC: Worldwide Business Analytics Software 2012–2016 Forecast and 2011 Vendor Shares

aws

AWS Big Data Portfolio

Collect


Amazon Kinesis Firehose


AWS Direct Connect


Amazon Kinesis Streams


Amazon Snowball

Store


Amazon S3


Amazon Glacier


Amazon Dynamo DB


Amazon RDS, Amazon Aurora


Amazon CloudSearch


Amazon Elasticsearch

Analyze


Amazon EMR


Amazon EC2


Amazon Redshift


Amazon Machine Learning


Amazon Kinesis Analytics


Amazon QuickSight


Amazon Athena


AWS Database Migration Service


AWS Data Pipeline


AWS Glue

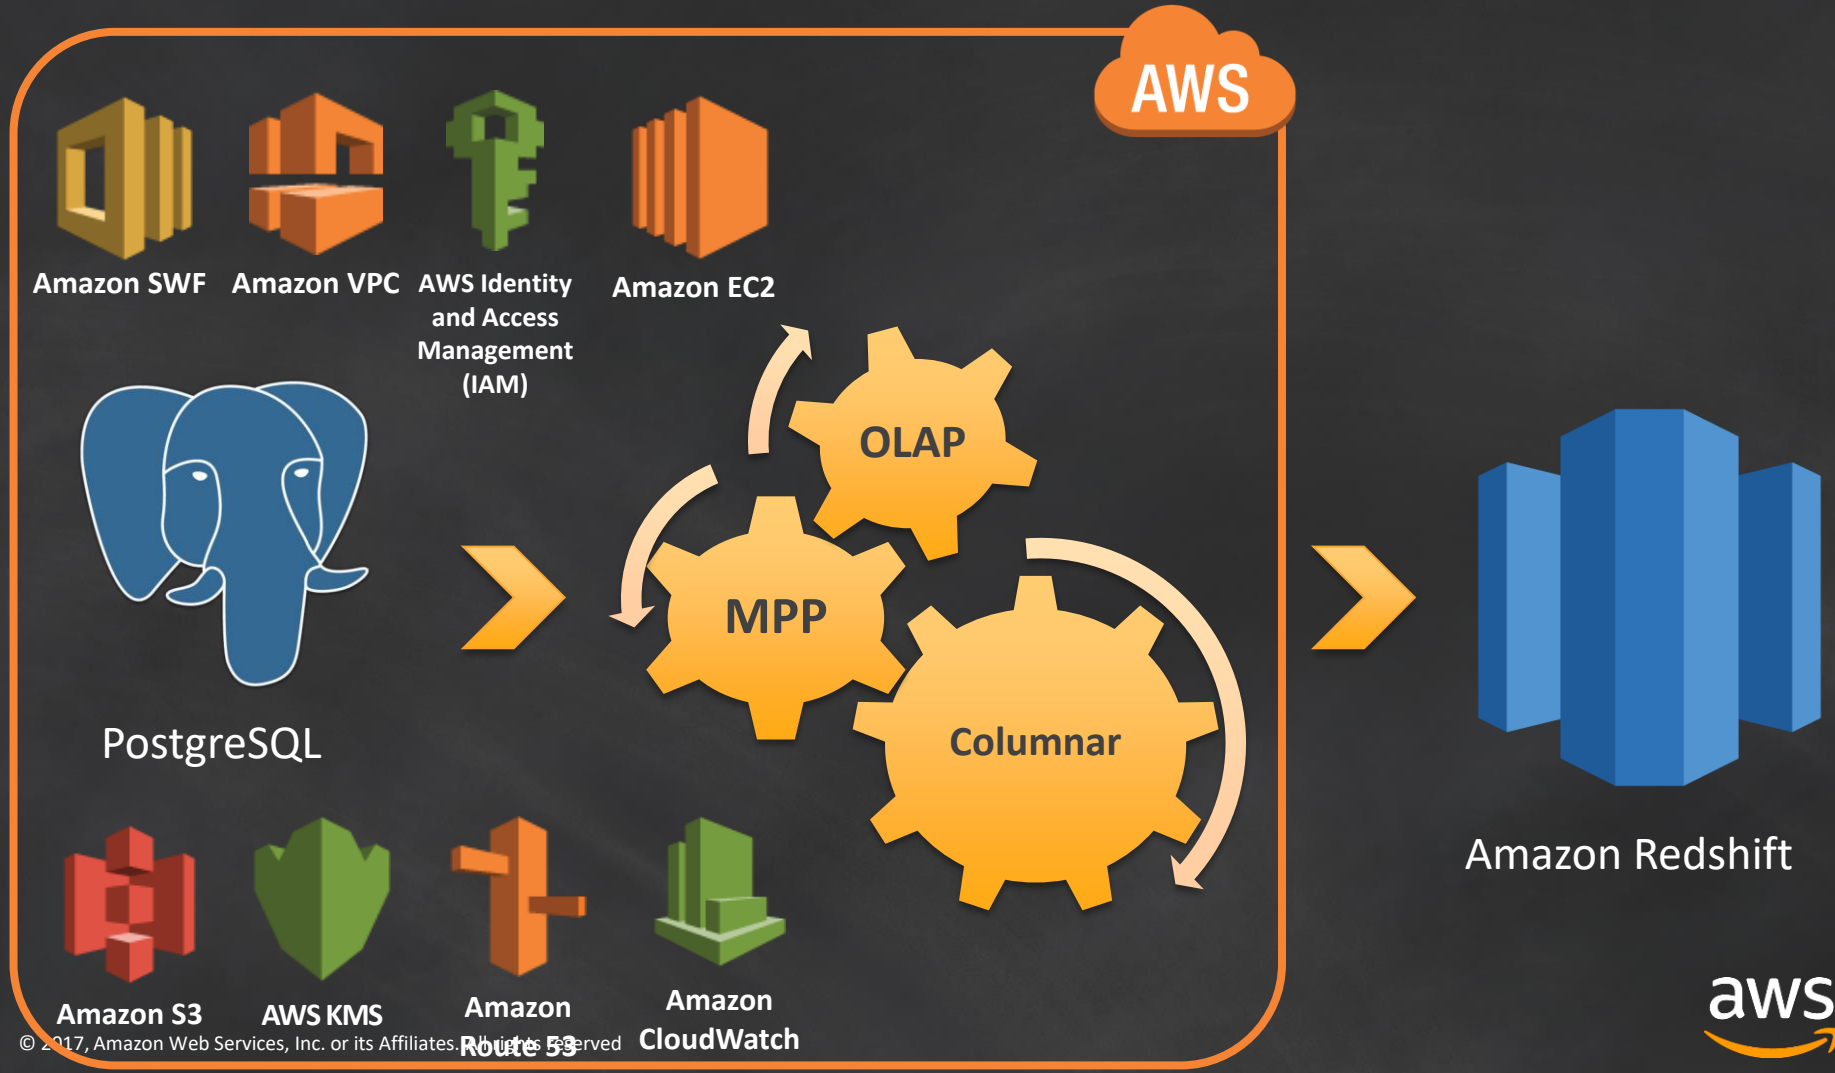

Amazon.com clickstream analytics

- Web log analysis for Amazon.com
  - PBs workload, 2TB/day@67% YoY
  - Largest table: 400 TB
- Understand customer behavior
- Previous solution
  - Legacy DW (Oracle)—query across 1 week/hr
  - Hadoop—query across 1 month/hr

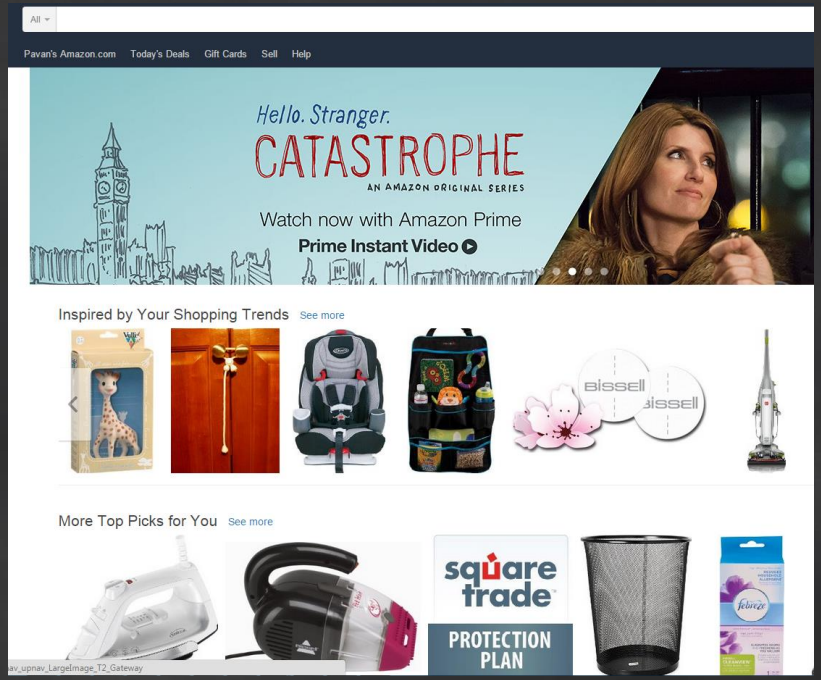

Results with Amazon Redshift

- Query 15 months **in 14 min**
- Load 5B rows in **10 min**
- 21B w/ 10B rows: **3 days to 2 hrs** (Hive → Redshift)
- Load pipeline: **90 hrs to 8 hrs** (Oracle → Redshift)

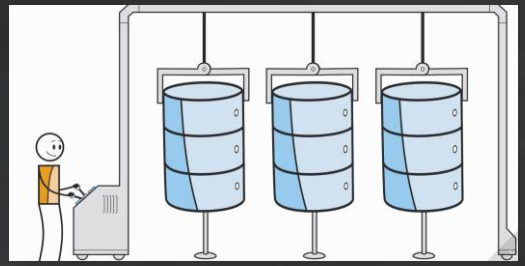

- 100 node DS2.8XL clusters
- Easy resizing
- Managed backups and restore
- Failure tolerance and recovery

- 20% time of one DBA
- Increased productivity

aws

Amazon Redshift

a lot faster
a lot simpler
a lot cheaper

Relational data warehouse

Massively parallel

Fully managed

HDD and SSD platforms

$1,000/TB/year; starts at $0.25/hour

aws

Selected Amazon Redshift Customers

Use Case: Traditional Data Warehousing

Business Reporting

Advanced pipelines and queries

Secure and Compliant

Bulk Loads and Updates

Easy Migration – Point & Click using AWS Database Migration Service

Secure & Compliant – End-to-End Encryption. SOC 1/2/3, PCI-DSS, HIPAA and FedRAMP compliant

Large Ecosystem – Variety of cloud and on-premises BI and ETL tools

docomo

Japanese Mobile Phone Provider

SCHOLASTIC

World's Largest Children's Book Publisher

Nasdaq

Powering 100 marketplaces in 50 countries

aws

Use Case: Log Analysis

Log & Machine IOT Data

Clickstream Events Data

Time-Series Data

Cheap – Analyze large volumes of data cost-effectively

Fast – Massively Parallel Processing (MPP) and columnar architecture for fast queries and parallel loads

Near real-time – Micro-batch loading and Amazon Kinesis Firehose for near-real time analytics

Pinterest

Interactive data analysis and recommendation engine

lyft

Ride analytics for pricing and product development

yelp

Ad prediction and on-demand analytics

aws

Use Case: Business Applications

Multi-Tenant BI Applications

Back-end services

Analytics as a Service

Fully Managed – Provisioning, backups, upgrades, security, compression all come built-in so you can focus on your business applications

Ease of Chargeback – Pay as you go, add clusters as needed. A few big common clusters & data marts

Service Oriented Architecture – Integrated with other AWS services. Easy to plug into your pipeline

Infosys

Infosys Information Platform (IIP)

accenture

Analytics-as-a-Service

Amplitude

Product and Consumer Analytics

aws

Amazon Redshift architecture

- **Leader node**
  - Simple SQL endpoint
  - Stores metadata
  - Optimizes query plan
  - Coordinates query execution
- **Compute nodes**
  - Local columnar storage
  - Parallel/distributed execution of all queries, loads, backups, restores, resizes
- **Start at just $0.25/hour, grow to 2 PB (compressed)**
  - DC1/DC2: SSD; scale from 160 GB to 326 TB
  - DS2: HDD; scale from 2 TB to 2 PB

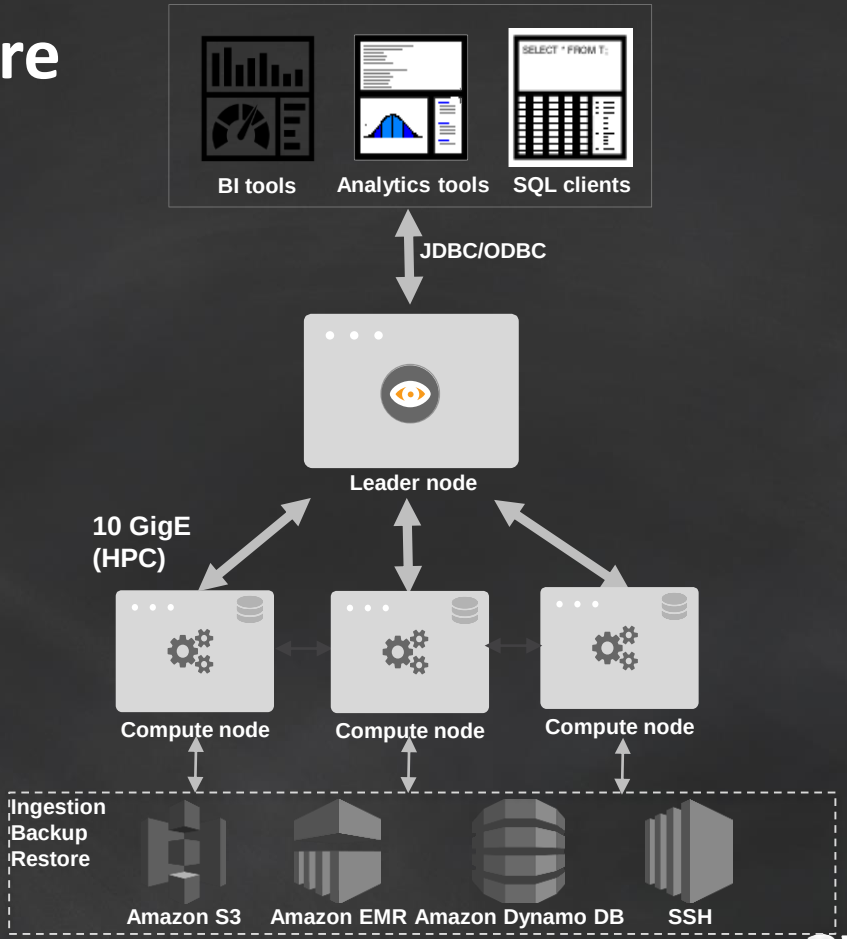

Benefit #1: Amazon Redshift is fast

- **Dramatically less I/O**
  - Column storage
  - Data compression
  - Zone maps
  - Direct-attached storage
  - Large data block sizes

```
analyze compression listing;

  Table   |      Column      | Encoding
----------+------------------+----------
 listing | listid           | delta
 listing | sellerid         | delta32k
 listing | eventid          | delta32k
 listing | dateid           | bytedict
 listing | numtickets       | bytedict
 listing | priceperticket   | delta32k
 listing | totalprice       | mostly32
 listing | listtime         | raw
```

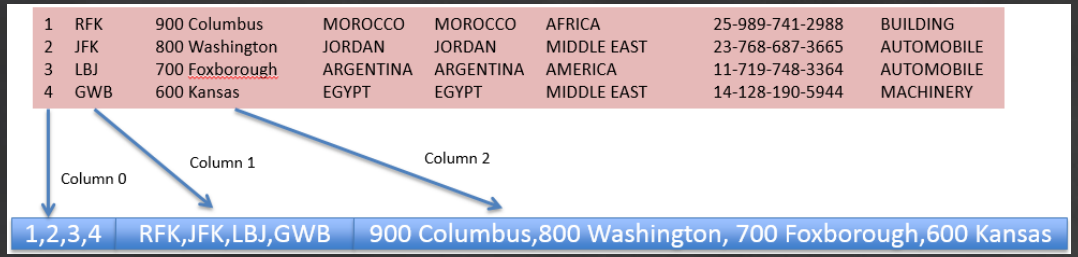

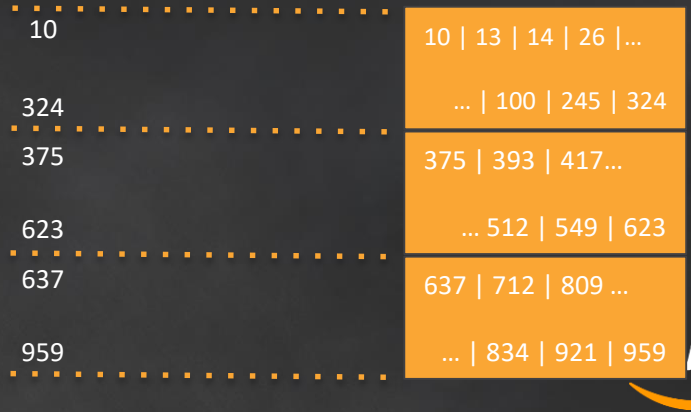

Benefit #2: Amazon Redshift is inexpensive

DS2 (HDD)	Price per hour for DS2.XL single node	Effective annual price per TB compressed
On-demand	$ 0.850	$ 3,725
1 year reservation	$ 0.500	$ 2,190
3 year reservation	$ 0.228	$ 999

DC1 (SSD)	Price per hour for DC1.L single node	Effective annual price per TB compressed
On-demand	$ 0.250	$ 13,690
1 year reservation	$ 0.161	$ 8,795
3 year reservation	$ 0.100	$ 5,500

Pricing is simple

Number of nodes x price/hour

No charge for leader node

No upfront costs

Pay as you go

aws

Benefit #3 : Amazon Redshift is easy to use

Provisioning in minutes

Automatic patching

SQL - Data loading

Backups are built-in

Security is built-in

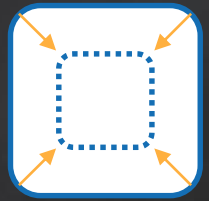

Compression is built-in

Benefit #4: Amazon Redshift is fully managed

Continuous/incremental backups

- Multiple copies within cluster
- Continuous and incremental backups to Amazon S3
- Continuous and incremental backups across regions
- Streaming restore

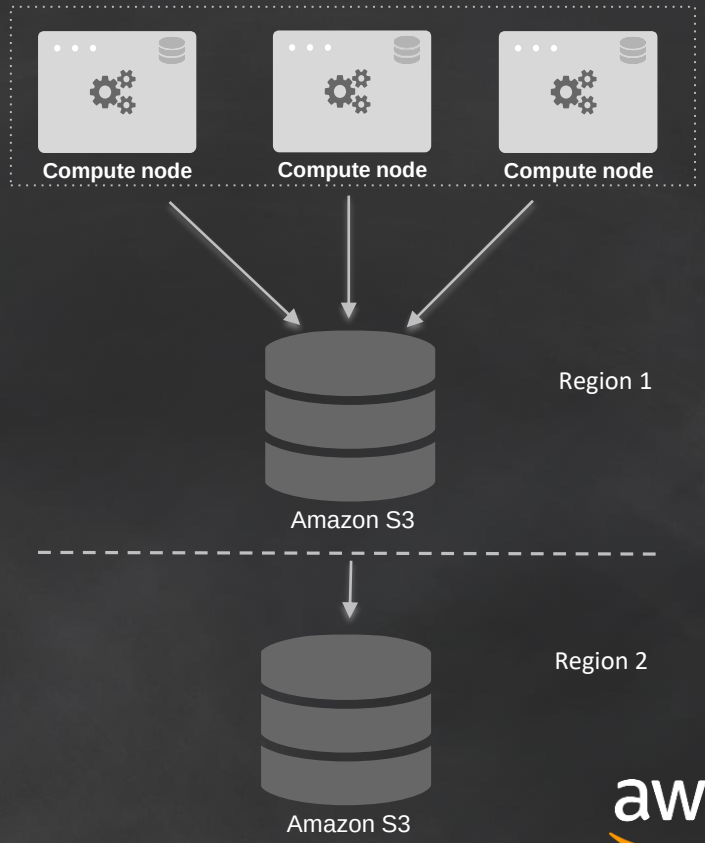

Benefit #5: Security is built-in

- Load encrypted from S3
- SSL to secure data in transit
  - ECDHE perfect forward secrecy
- Amazon VPC for network isolation
- Encryption to secure data at rest
  - All blocks on disks and in S3 encrypted
  - Block key, cluster key, master key (AES-256)
  - On-premises HSM & AWS CloudHSM support
- Audit logging and AWS CloudTrail integration
- SOC 1/2/3, PCI-DSS, FedRAMP, BAA

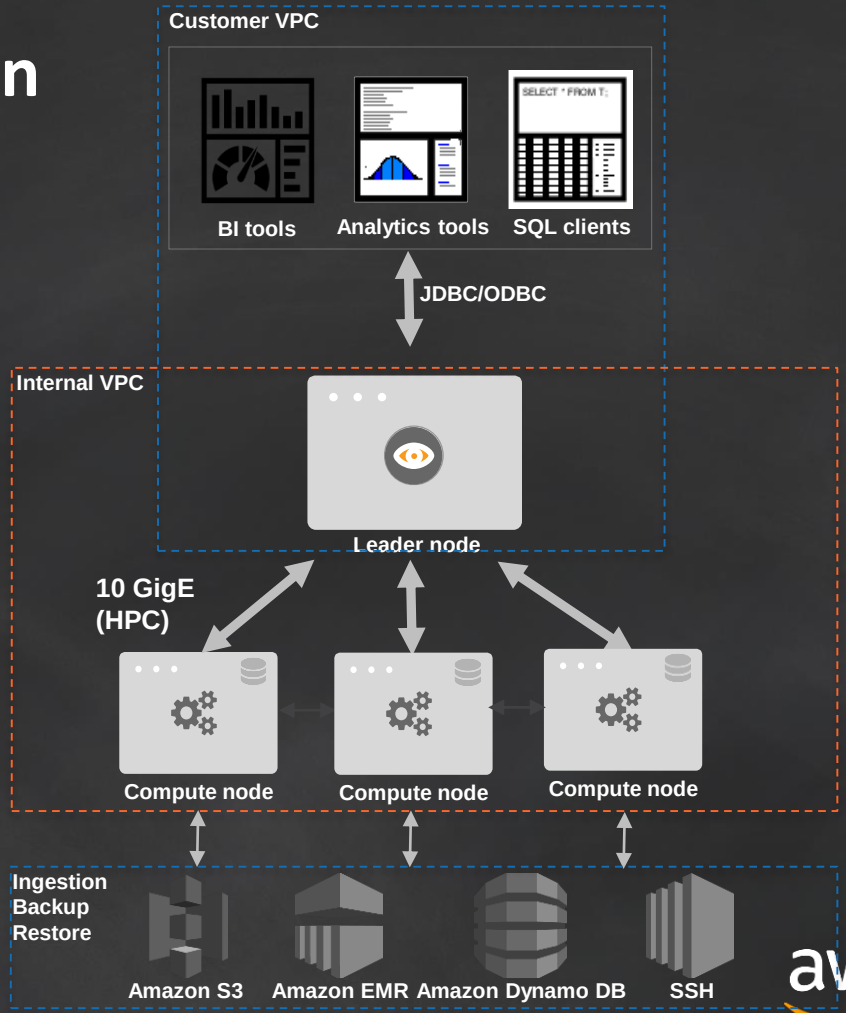

Benefit #6: Amazon Redshift is powerful

- Approximate functions
- User defined functions
- Machine learning
- Data science

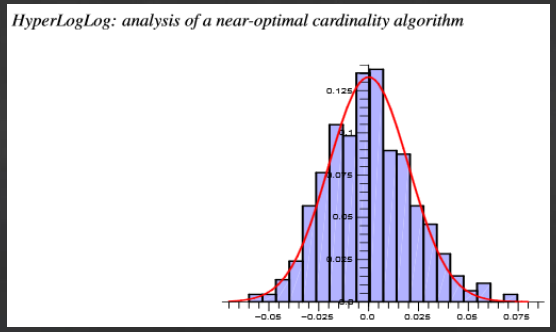

Benefit #7: Amazon Redshift has a large ecosystem

aws

Amazon Redshift Spectrum

Amazon Redshift Spectrum

Run SQL queries directly against data in S3 using thousands of nodes

Fast @ exabyte scale

Elastic & highly available

On-demand, pay-per-query

High concurrency: Multiple clusters access same data

No ETL: Query data in-place using open file formats

Full Amazon Redshift SQL support

aws

Extend Redshift Queries To Amazon S3

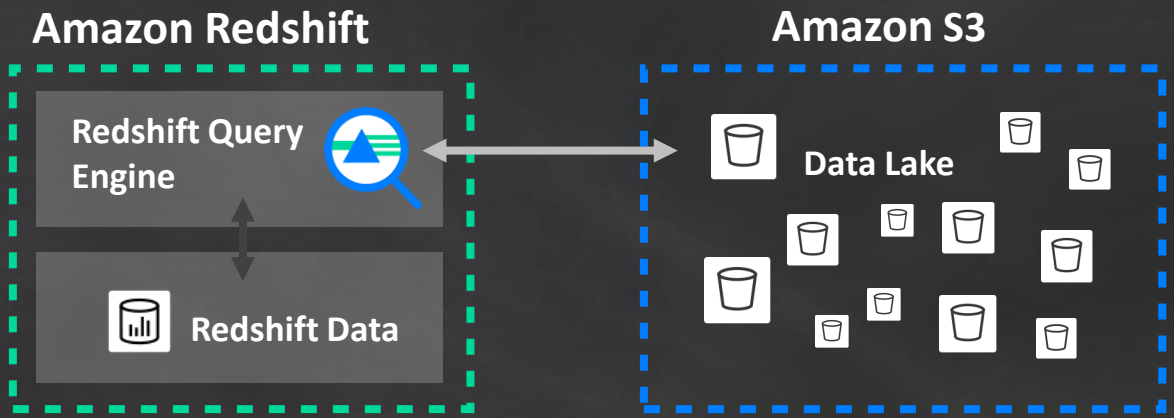

Query S3 directly or join data across Redshift and S3

Scale Redshift compute and storage separately

Support for CSV, Parquet, ORC, Grok, Avro and more formats

Life of a query

1 Query
SELECT COUNT(*)
FROM S3.EXT_TABLE
GROUP BY…

JDBC/ODBC

Amazon Redshift

2 Query is optimized and compiled at the leader node. Determine what gets run locally and what goes to Amazon Redshift Spectrum

3 Query plan is sent to all compute nodes

4 Compute nodes obtain partition info from Data Catalog; dynamically prune partitions

5 Each compute node issues multiple requests to the Amazon Redshift Spectrum layer

6 Amazon Redshift Spectrum nodes scan your S3 data

7 Amazon Redshift Spectrum projects, filters, joins and aggregates

8 Final aggregations and joins with local Amazon Redshift tables done in-cluster

9 Result is sent back to client

1 2 3 4 … N

Amazon S3
Exabyte-scale object storage

Data Catalog
Glue / Apache Hive Metastore

aws

Use cases

- Extend data warehousing to data lake in Amazon S3
- Cost reduction for less frequently accessed data
- Better query performance and minimal time to insight at large scale
- Query many open formats and large data sets that can’t be loaded into cluster
- Improved concurrency with multiple Redshift clusters querying common data
- Elastically scale compute resources separately from the storage layer in S3
- Use familiar SQL to cleanse, transform and load data from S3 to Redshift

Multiple cluster architecture

- Use multiple Redshift clusters to query same copy of data in s3

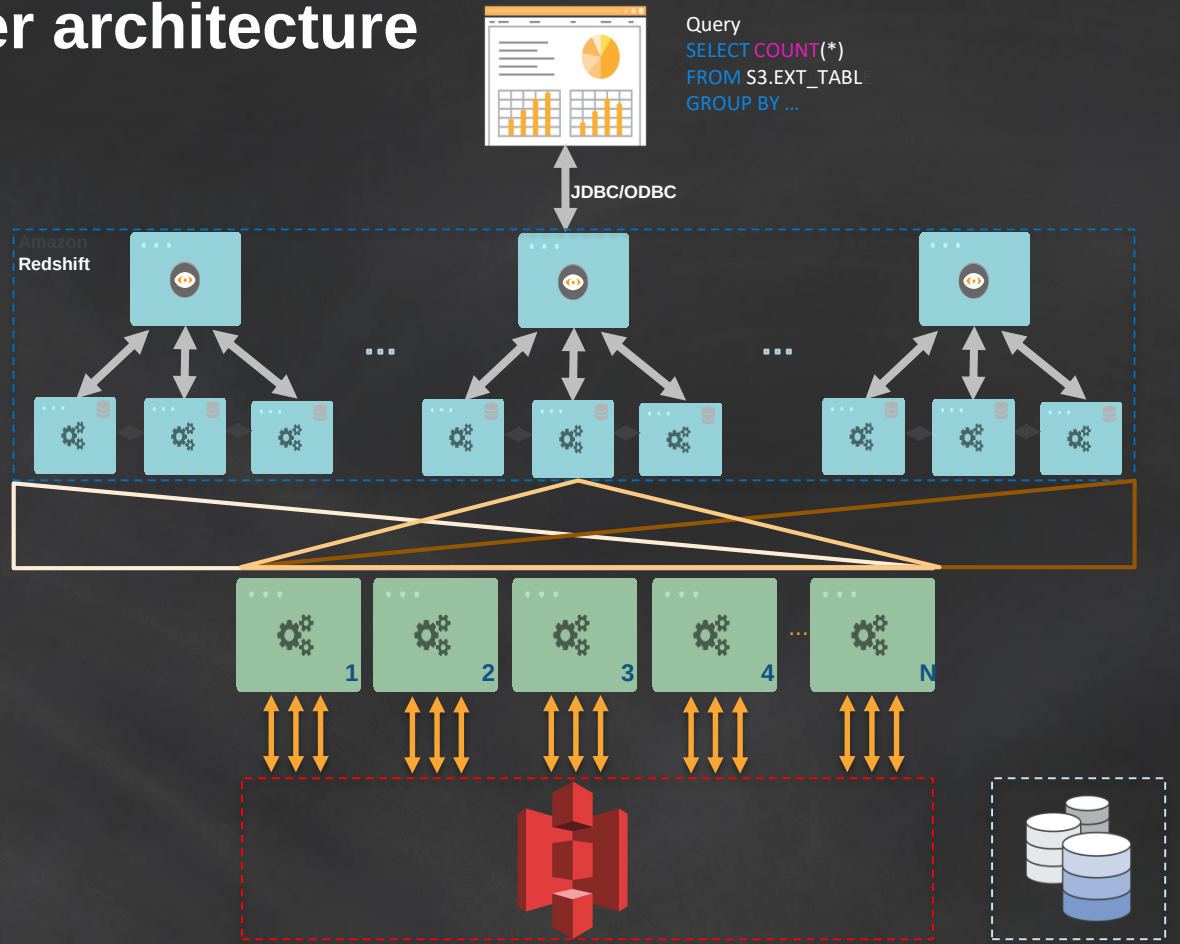

Simplify and accelerate your ETL pipelines

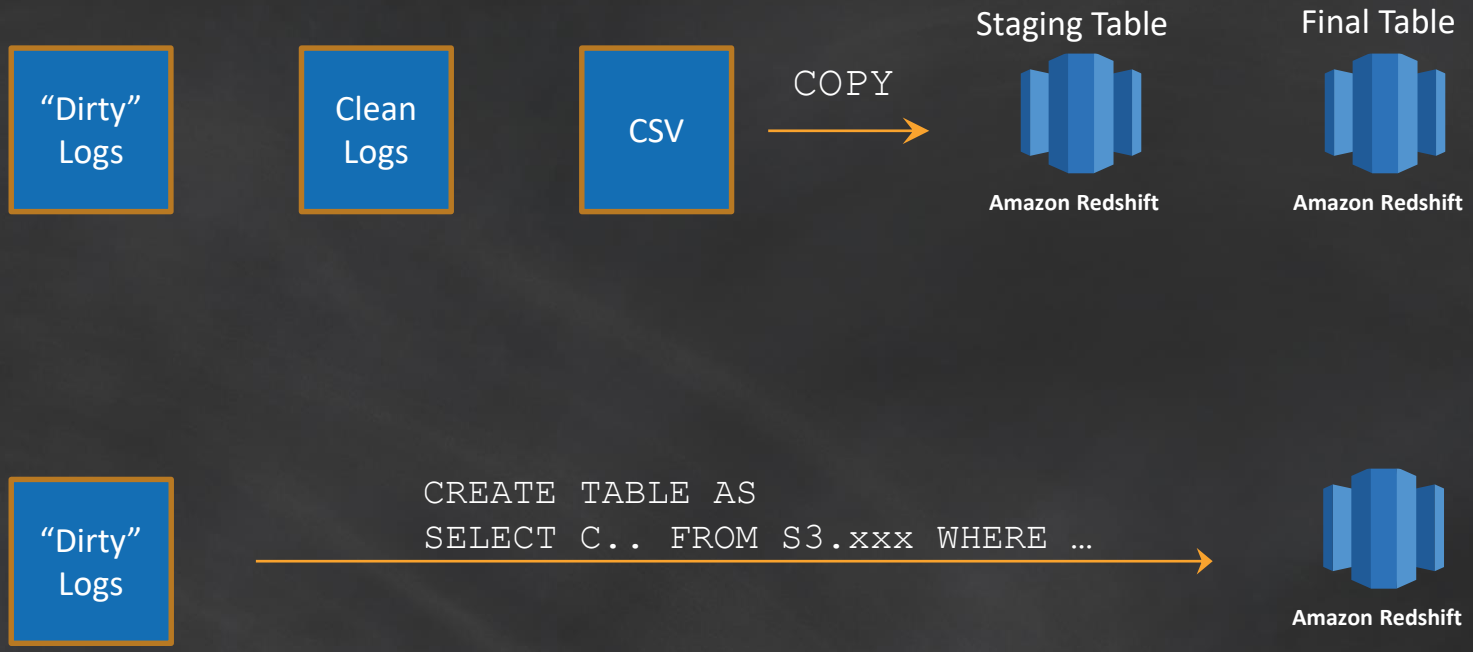

Customer Use Cases

Nasdaq: powering 100 marketplaces in 50 countries

Orders, quotes, trade executions, market “tick” data from 7 exchanges

7 billion rows/day

Analyze market share, client activity, surveillance, billing, and so on

Microsoft SQL Server on-premises

Expensive legacy DW ($1.16 M/yr.)

Limited capacity (1 yr. of data online)

Needed lower TCO

Must satisfy multiple security and regulatory requirements

Similar performance

aws

Nasdaq: powering 100 marketplaces in 50 countries

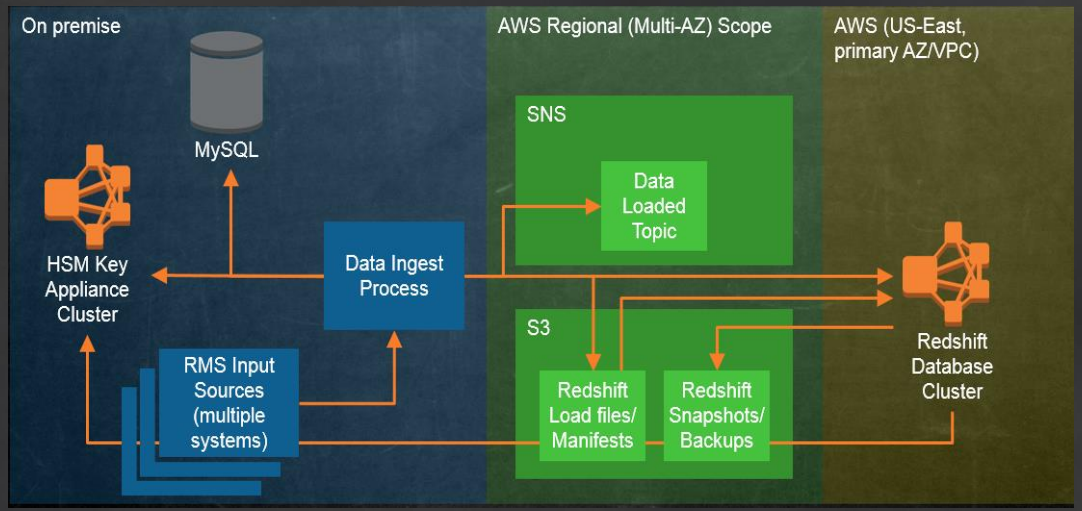

23 node DS2.8XL cluster

828 vCPUs, 5 TB RAM

368 TB compressed

2.7 T rows, 900 B derived

8 tables with 100 B rows

7 man-month migration

¼ the cost, 2x storage, room to grow

Faster performance, very secure

Customers love Amazon Redshift Spectrum

Time Inc.

"Redshift Spectrum enables us to directly operate on our data in its native format in Amazon S3 with no preprocessing or transformation."

NTT docomo

"Redshift Spectrum will let us expand the universe of the data we analyze to 100s of petabytes over time. This is truly a game changer, and we can think of no other system in the world that can get us there."

edmunds

"Redshift Spectrum's fast performance across massive data sets is unprecedented."

REDFIN

"Our data science team using Amazon EMR can now collaborate with our marketing and product teams using Redshift Spectrum to analyze the same Amazon S3 data sets."

yelp

"Multiple teams can now query the same Amazon S3 data sets using both Redshift and EMR."

RECRUIT
Recruit Technologies Co.,Ltd.

"Redshift Spectrum will help us scale yet further while also lowering our costs."

aws

Getting Started

Provisioning

Enter cluster details

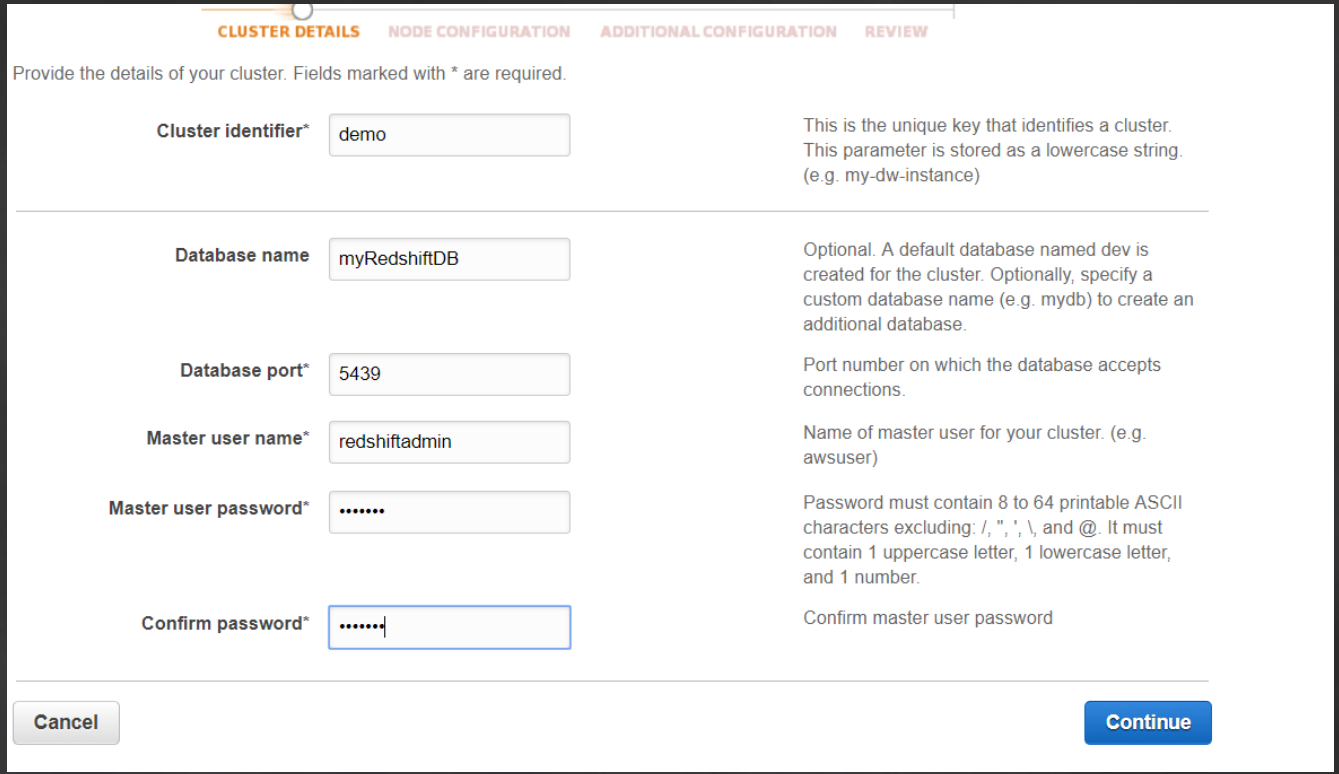

Select node configuration

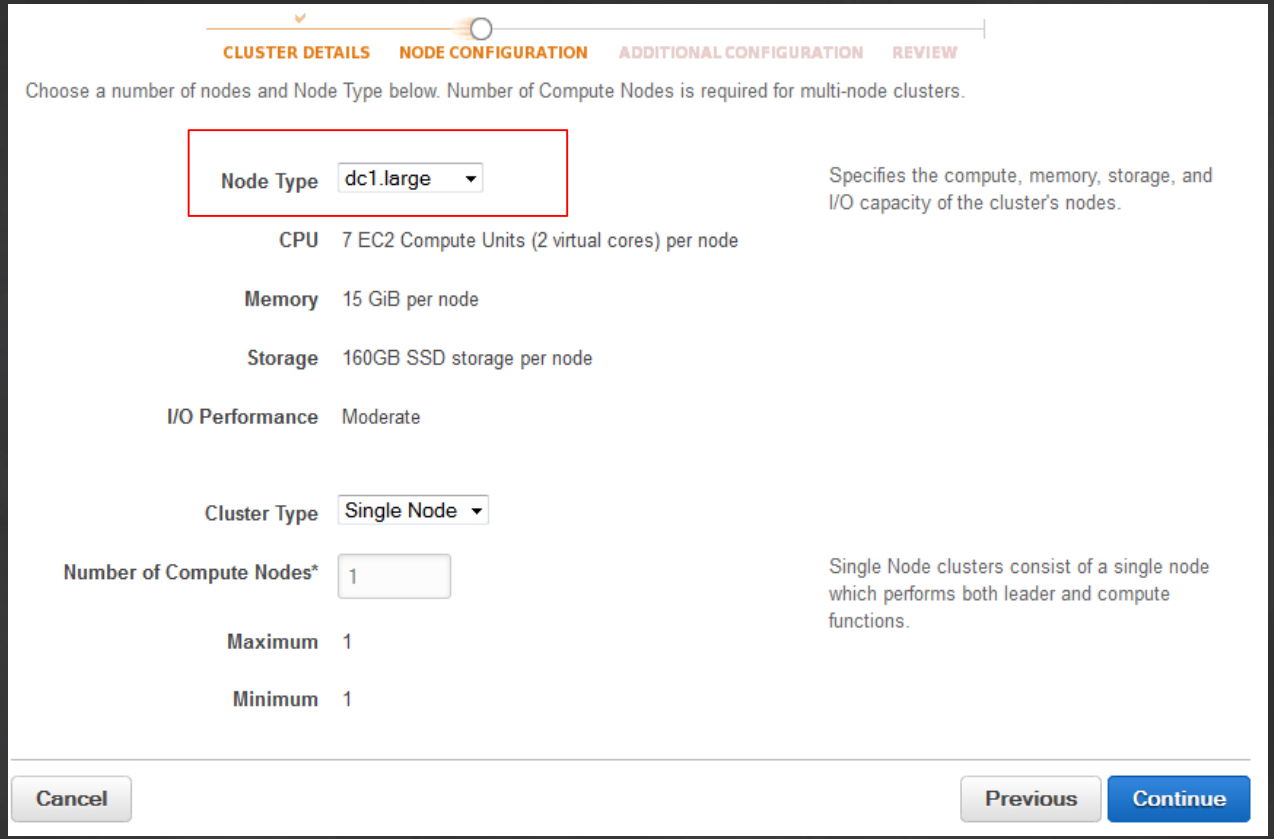

Select security settings and provision

Point-and-click resize

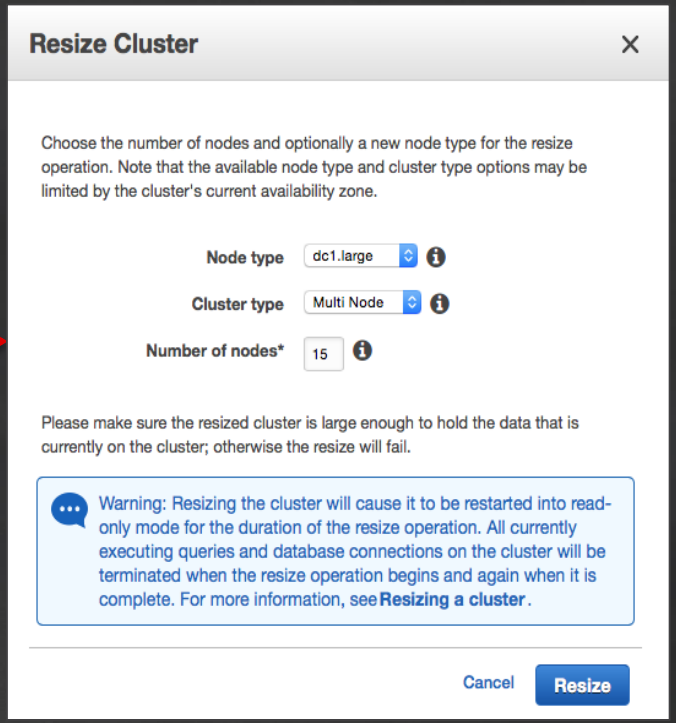

Data Modeling

Sort Keys

Zone maps

SELECT COUNT(*) FROM LOGS WHERE DATE = '09-JUNE-2013'

Unsorted table

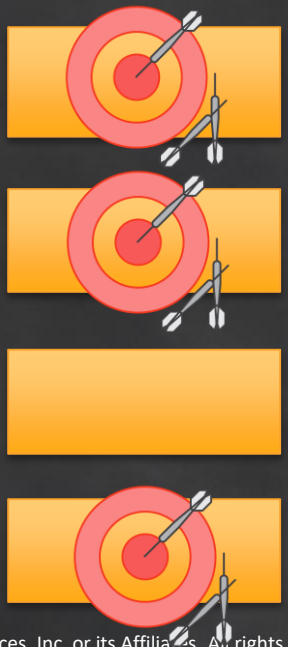

MIN: 01-JUNE-2013
MAX: 20-JUNE-2013

MIN: 08-JUNE-2013
MAX: 30-JUNE-2013

MIN: 12-JUNE-2013
MAX: 20-JUNE-2013

MIN: 02-JUNE-2013
MAX: 25-JUNE-2013

Sorted by date

MIN: 01-JUNE-2013
MAX: 06-JUNE-2013

MIN: 07-JUNE-2013
MAX: 12-JUNE-2013

MIN: 13-JUNE-2013
MAX: 18-JUNE-2013

MIN: 19-JUNE-2013
MAX: 24-JUNE-2013

aws

Sort Keys

- Single column
- Compound
- Interleaved

Sort Keys

Single Column

- Table is sorted by 1 column

```
[ SORTKEY ( date ) ]
```

Date	Region	Country
2-JUN-2015	Oceania	New Zealand
2-JUN-2015	Asia	Singapore
2-JUN-2015	Africa	Zaire
2-JUN-2015	Asia	Hong Kong
3-JUN-2015	Europe	Germany
3-JUN-2015	Asia	Korea

- Best for:
  - Queries that use 1st column (i.e. *date)* as primary filter
  - Can speed up joins and group bys
  - Quickest to VACUUM

aws

Sort Keys

Compound

- Table is sorted by 1st column , then 2nd column etc.

```
[ SORTKEY COMPOUND ( date, region, country) ]
```

Date	Region	Country
2-JUN-2015	Africa	Zaire
2-JUN-2015	Asia	Korea
2-JUN-2015	Asia	Singapore
2-JUN-2015	Europe	Germany
3-JUN-2015	Asia	Hong Kong
3-JUN-2015	Asia	Korea

- Best for:
  - Queries that use 1st column as primary filter, then other cols
  - Can speed up joins and group bys
  - Slower to VACUUM

aws

Distribution

- EVEN
- KEY
- ALL

Distribution

ID	Gender	Name
101	M	John Smith
292	F	Jane Jones
139	M	Peter Black
446	M	Pat Partridge
658	F	Sarah Cyan
164	M	Brian Snail
209	M	James White
306	F	Lisa Green

EVEN

ID	Gender	Name
101	M	John Smith
306	F	Lisa Green

ID	Gender	Name
292	F	Jane Jones
209	M	James White

ID	Gender	Name
139	M	Peter Black
164	M	Brian Snail

ID	Gender	Name
446	M	Pat Partridge
658	F	Sarah Cyan

DISTSTYLE EVEN

Distribution

ID	Gender	Name
101	M	John Smith
292	F	Jane Jones
139	M	Peter Black
446	M	Pat Partridge
658	F	Sarah Cyan
164	M	Brian Snail
209	M	James White
306	F	Lisa Green

KEY

1

ID	Gender	Name
101	M	John Smith
306	F	Lisa Green

2

ID	Gender	Name
292	F	Jane Jones
209	M	James White

3

ID	Gender	Name
139	M	Peter Black
164	M	Brian Snail

4

ID	Gender	Name
446	M	Pat Partridge
658	F	Sarah Cyan

DISTSTYLE KEY

Distribution

ID	Gender	Name
101	M	John Smith
292	F	Jane Jones
139	M	Peter Black
446	M	Pat Partridge
658	F	Sarah Cyan
164	M	Brian Snail
209	M	James White
306	F	Lisa Green

ID	Gender	Name
101	M	John Smith
139	M	Peter Black
446	M	Pat Partridge
164	M	Brian Snail
209	M	James White

ID	Gender	Name
292	F	Jane Jones
658	F	Sarah Cyan
306	F	Lisa Green

DISTSTYLE KEY

Distribution

ID	Gender	Name
101	M	John Smith
292	F	Jane Jones
139	M	Peter Black
446	M	Pat Partridge
658	F	Sarah Cyan
164	M	Brian Snail
209	M	James White
306	F	Lisa Green

ALL

101	M	John Smith
292	F	Jane Jones
139	M	Peter Black
446	M	Pat Partridge
658	F	Sarah Cyan
164	M	Brian Snail
209	M	Lisa Green
306	F	James White

101	M	John Smith
292	F	Jane Jones
139	M	Peter Black
446	M	Pat Partridge
658	F	Sarah Cyan
164	M	Brian Snail
209	M	Lisa Green
306	F	James White

101	M	John Smith
292	F	Jane Jones
139	M	Peter Black
446	M	Pat Partridge
658	F	Sarah Cyan
164	M	Brian Snail
209	M	Lisa Green
306	F	James White

101	M	John Smith
292	F	Jane Jones
139	M	Peter Black
446	M	Pat Partridge
658	F	Sarah Cyan
164	M	Brian Snail
209	M	Lisa Green
306	F	James White

DISTSTYLE ALL

Distribution

- EVEN
  - Tables with no joins or group by
- KEY
  - Large Fact tables
  - Large dimension tables
- ALL
  - Medium dimension tables (1K – 2M)
  - Small dimension tables

aws

Loading Data

Data loading options

Data loading options

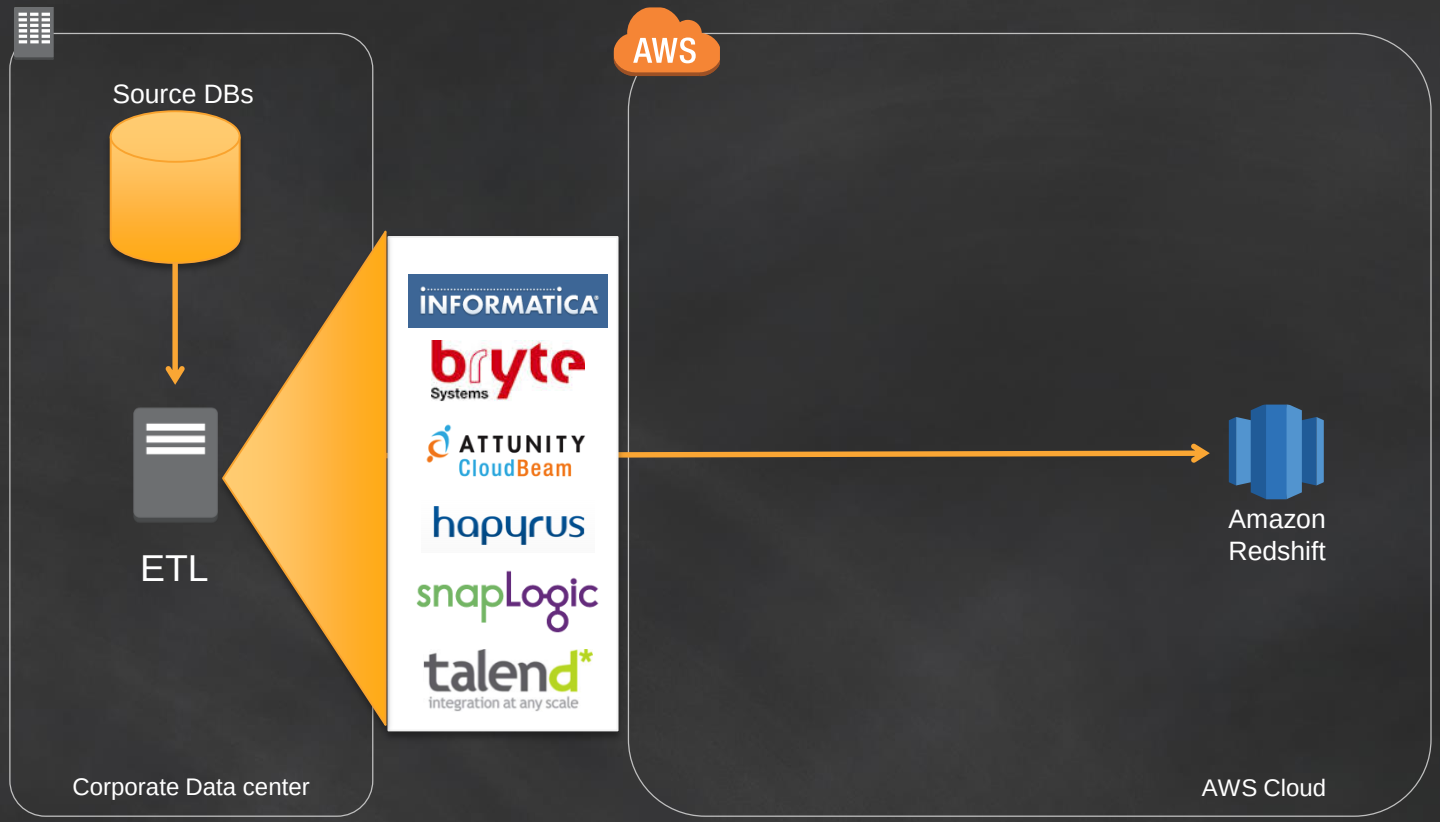

Data loading options

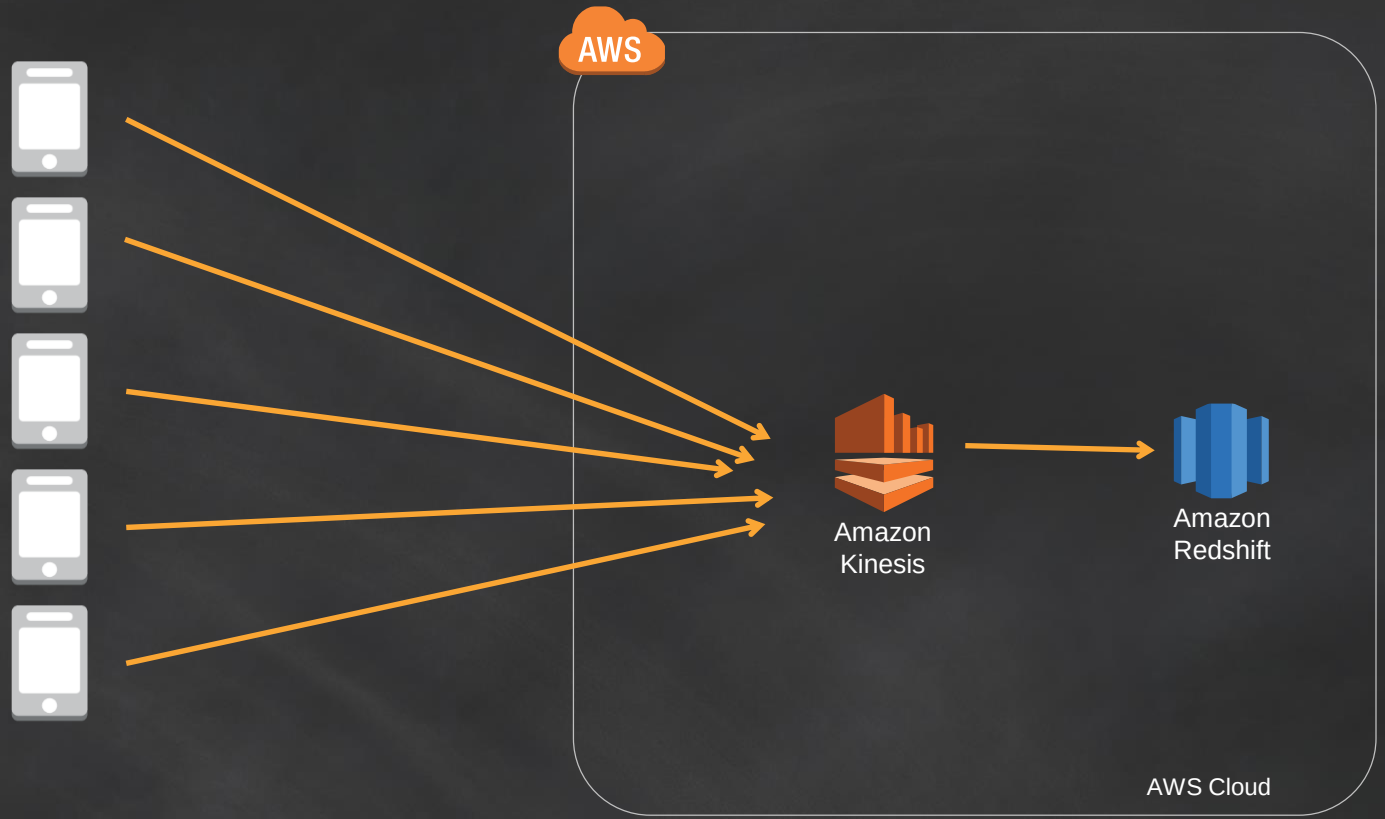

ETL data into your data warehouse

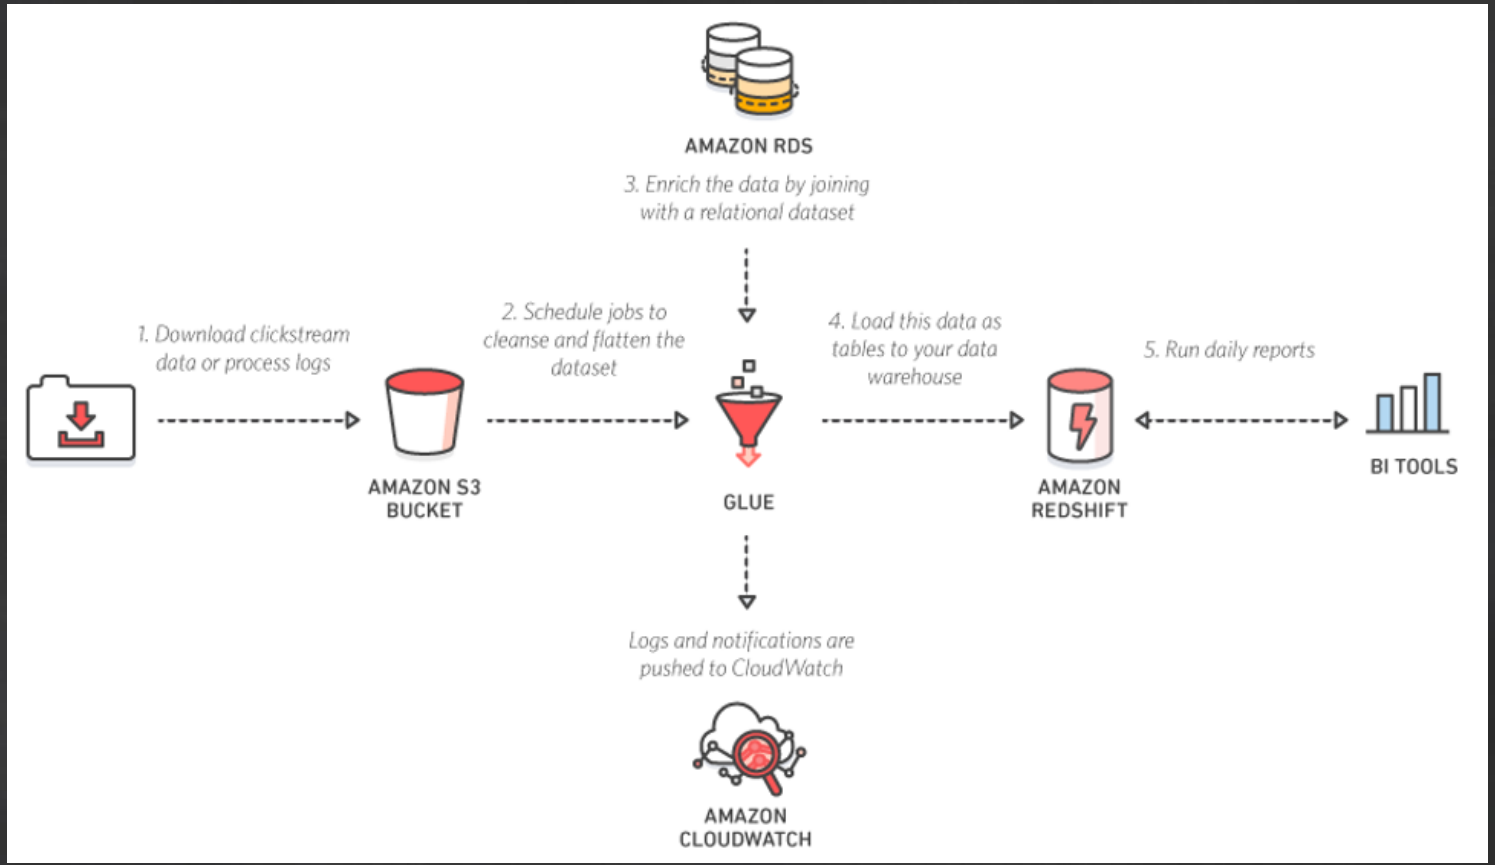

Instantly query your data lake on Amazon S3

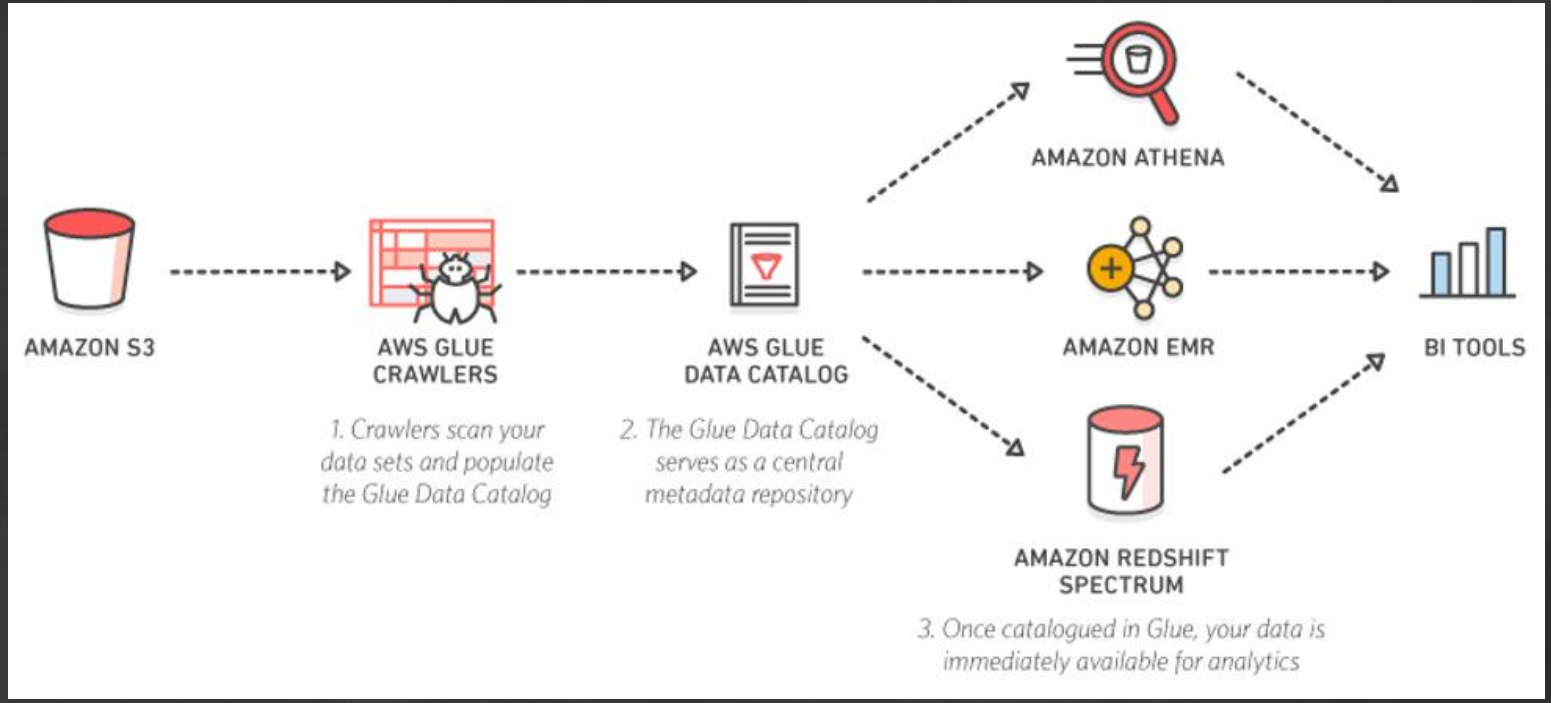

Redshift Spectrum: Defining External Schema and External Tables

1. Define an external schema in Amazon Redshift using the AWS Glue Data Catalog or your own Apache Hive Metastore

```
CREATE EXTERNAL SCHEMA <schema_name>
FROM { [ DATA CATALOG ] | HIVE METASTORE }
DATABASE 'database_name'
IAM_ROLE 'iam-role-arn'
```

2. Register external tables using Amazon Athena, your Hive Metastore client, or from Amazon Redshift

```
CREATE EXTERNAL TABLE <schema_name>.<table_name>
[PARTITIONED BY <column_name, data_type, …>]
STORED AS file_format
LOCATION s3_location
[TABLE PROPERTIES property_name=property_value, …];
```

3. Query external tables

```
SELECT … FROM <schema_name>.<table_name> …
```


aws

Querying

Amazon Redshift works with your existing BI tools

Amazon Redshift

JDBC/ODBC

actian · pentaho · ACTUATE The BIRT Company · periscope · birst · Qlik · ClearStory DATA Now You See It · SiSense · Chart.io Turn data to insight · tableau SOFTWARE · Logi ANALYTICS · TIBCO · looker · Yellowfin · MicroStrategy · MODE

aws

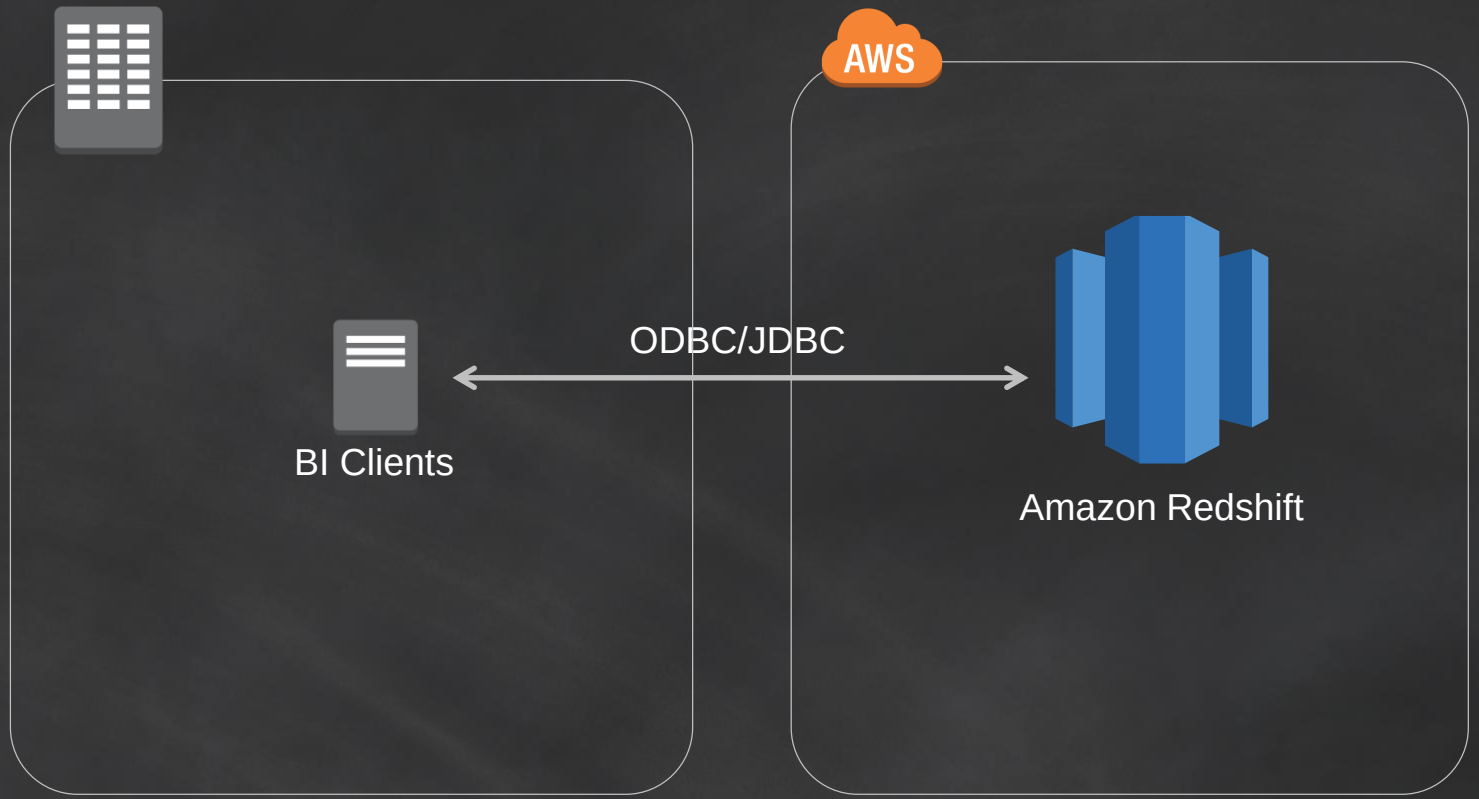

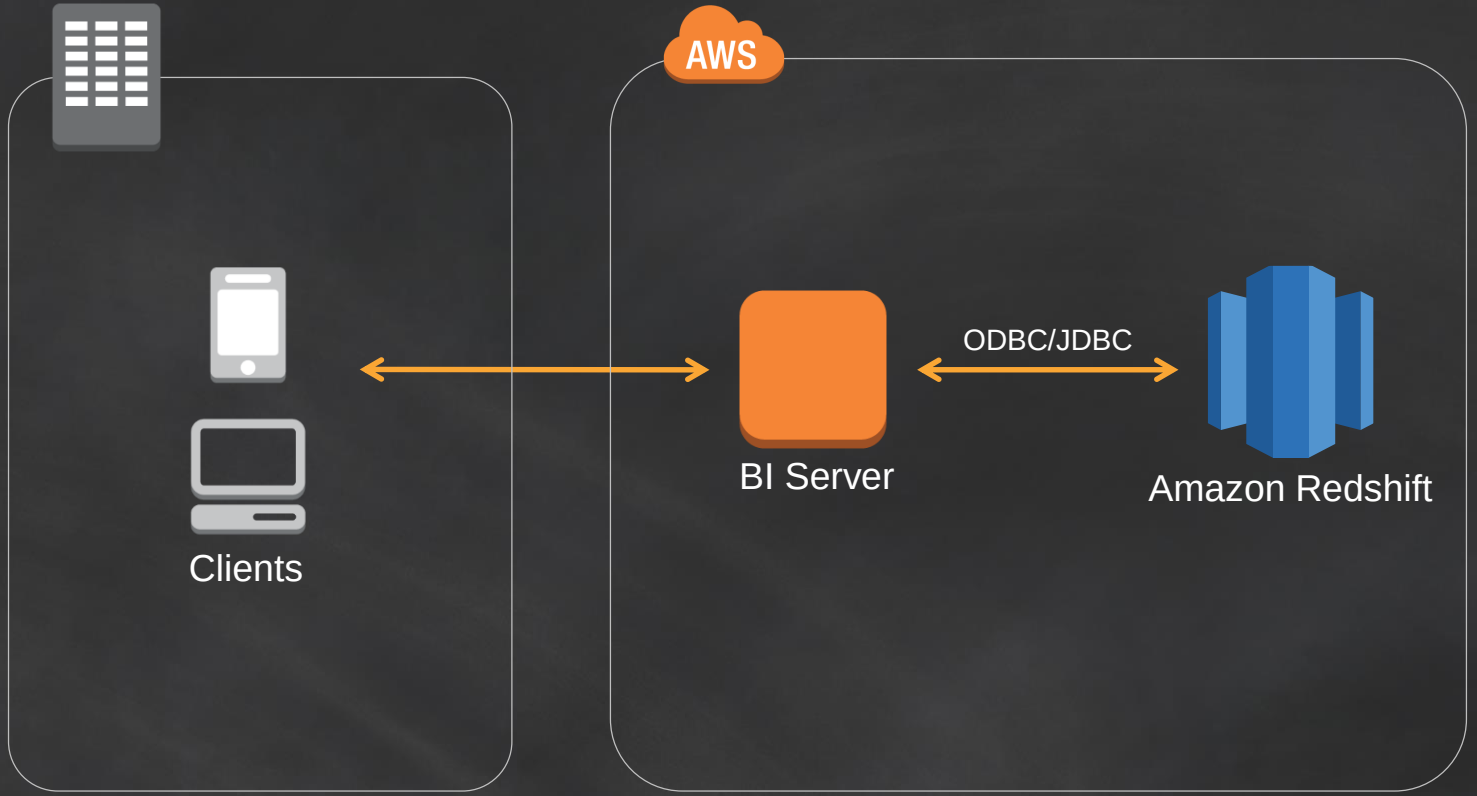

View explain plans

Query Execution Details

Plan | Actual

```
XN Hash Join DS_BCAST_INNER  (cost=52602972108.91..68656434844.99 rows=23608 width=71)
  ->  XN Seq Scan on supplier s  (cost=0.00..10000.00 rows=1000000 width=22)
  ->  XN Hash  (cost=52602972049.89..52602972049.89 rows=23608 width=57)
        ->  XN Hash Join DS_BCAST_INNER  (cost=1752.83..52602972049.89 rows=23608 width=57)
              ->  XN Hash Join DS_BCAST_INNER  (cost=1718.72..52499650202.87 rows=70082 width=61)
              ->  XN Hash  (cost=31.95..31.95 rows=861 width=4)
                    ->  XN Hash Join DS_BCAST_INNER  (cost=624.99..14000065191.80 rows=714354 wid
                    ->  XN Hash  (cost=874.99..874.99 rows=87499 width=33)
                          ->  XN Seq Scan on lineorder lo  (cost=0.00..149965.90 rows=14996590 wi
                          ->  XN Hash  (cost=499.99..499.99 rows=49999 width=20)
                                ->  XN Seq Scan on part p  (cost=0.00..499.99 rows=49999 width=20
```

Plan | Actual

Expand Level | Collapse Level | Expand All

- Hash Join (Broadcasted Inner Table)
  - Sequential Scan on supplier s
  - Hash
    - Hash Join (Broadcasted Inner Table)
      - Hash Join (Broadcasted Inner Table
        - Hash Join (Broadcasted Inner Tab
          - Sequential Scan on lineorder lo
          - Hash
            - Sequential Scan on part p
        - Hash
          - Sequential Scan on customer c
      - Hash
        - Sequential Scan on dwdate d

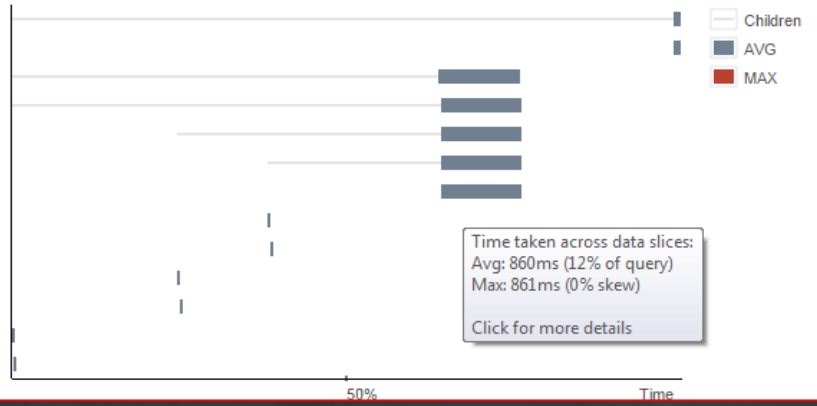


```
select
    lo_orderkey,
    p_name,
    c_name,
    s_address,
    lo_quantity
from
    lineorder lo,
    part p,
    supplier s,
    customer c,
    dwdate d
where
    lo_custkey = c_custkey
    and     lo_partkey = p_partkey
    and     lo_suppkey = s_suppkey
    and     lo_orderdate = d_datekey
    and     d_sellingseason = 'Summer'
```

aws

Monitor query performance

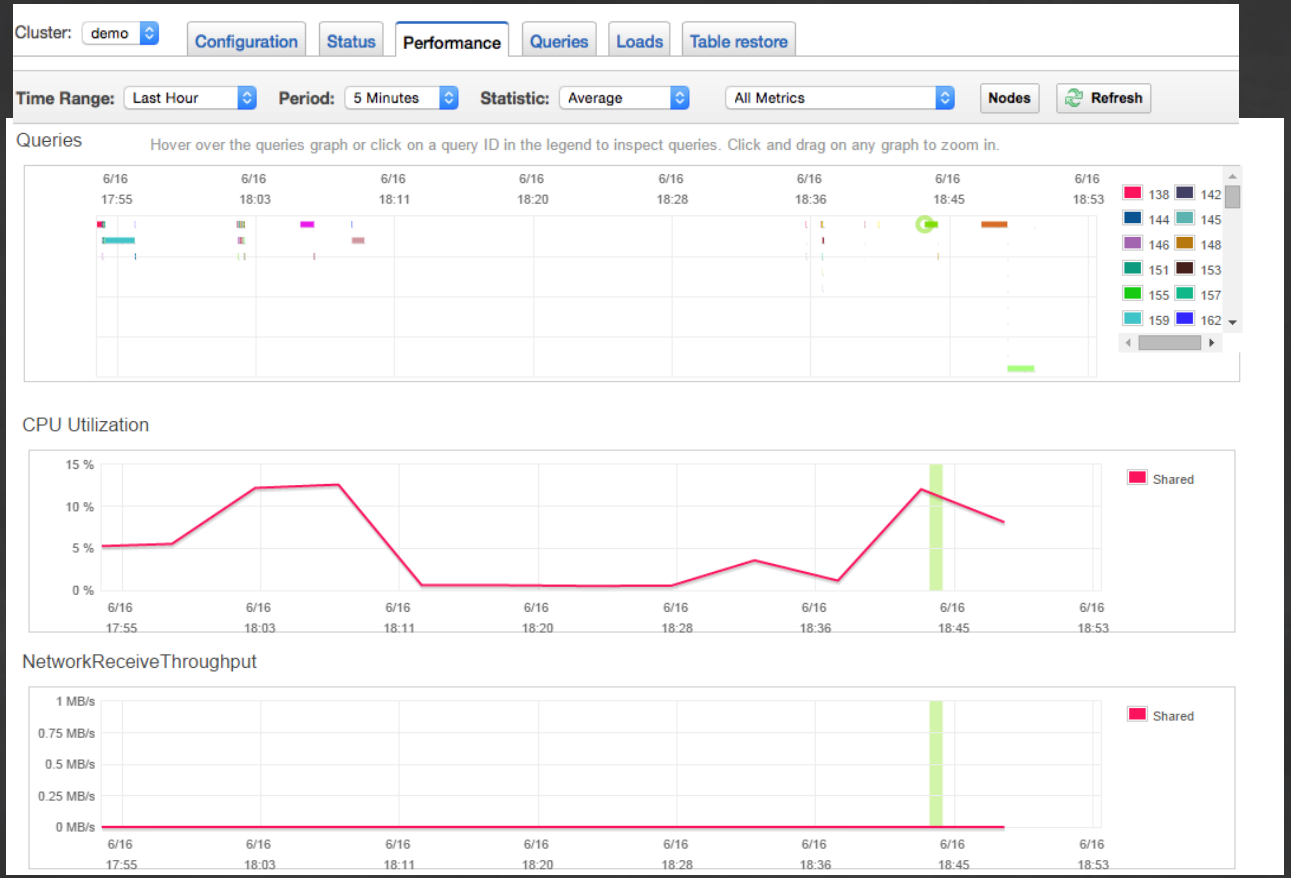

Resources

- **Detail Pages**
  - http://aws.amazon.com/redshift
  - https://aws.amazon.com/redshift/spectrum
  - https://aws.amazon.com/marketplace/redshift/
  - https://aws.amazon.com/redshift/developer-resources/
  - Amazon Redshift Utilities - GitHub

- **Best Practices**
  - http://docs.aws.amazon.com/redshift/latest/dg/c_loading-data-best-practices.html
  - http://docs.aws.amazon.com/redshift/latest/dg/c_designing-tables-best-practices.html
  - http://docs.aws.amazon.com/redshift/latest/dg/c-optimizing-query-performance.html

Everything and Anything Startups Need to Get Started on AWS

aws.amazon.com/activate