



Best Practices for Migrating your Data Warehouse to Amazon Redshift


Darin Briskman, AWS Technical Evangelist
briskman@amazon.com

Why Migrate to Amazon Redshift?



Transactional database

- Redshift can be 100x faster
- Scales from GBs to PBs
- Analyze data without storage constraints



MPP database

- Redshift is 10x cheaper
- Easy to provision and operate
- Higher productivity



Hadoop

- Redshift can be 10x faster
- No programming
- Standard interfaces and integration to leverage BI tools, machine learning, streaming

Migration from Oracle @ Boingo Wireless

2000+ Commercial Wi-Fi locations

1 million+ Hotspots

90M+ ad engagements

100+ countries



Legacy DW: Oracle 11g based DW

Before migration

Rapid data growth slowed analytics

Mediocre IOPS, limited memory, vertical scaling

Admin overhead

Expensive (license, h/w, support)

After migration


180x performance improvement

7x cost savings


Migration from Oracle @ Boingo Wireless



7X cheaper than Oracle Exadata



180X faster than Oracle database



Migration from Greenplum @ NTT Docomo

68 million customers

10s of TBs per day of data across mobile network

6PB of total data (uncompressed)

Data science for marketing operations, logistics etc.

Legacy DW: Greenplum on-premises

After migration:

125 node DS2.8XL cluster

4,500 vCPUs, 30TB RAM

6 PB uncompressed

10x faster analytic queries

50% reduction in time for new BI app. deployment

Significantly less ops. overhead



Migration from SQL on Hadoop @ Yahoo

- Analytics for website/mobile events across multiple Yahoo properties
- On an average day
 - 2B events
 - 25M devices

Before migration: Hive – Found it to be slow, hard to use, share and repeat

After migration:

21 node DC1.8XL (SSD)


50TB compressed data

100x performance improvement


Real-time insights

Easier deployment and maintenance

Migration from SQL on Hadoop @ Yahoo



Amazon Redshift Migration Process



How to Migrate?

3

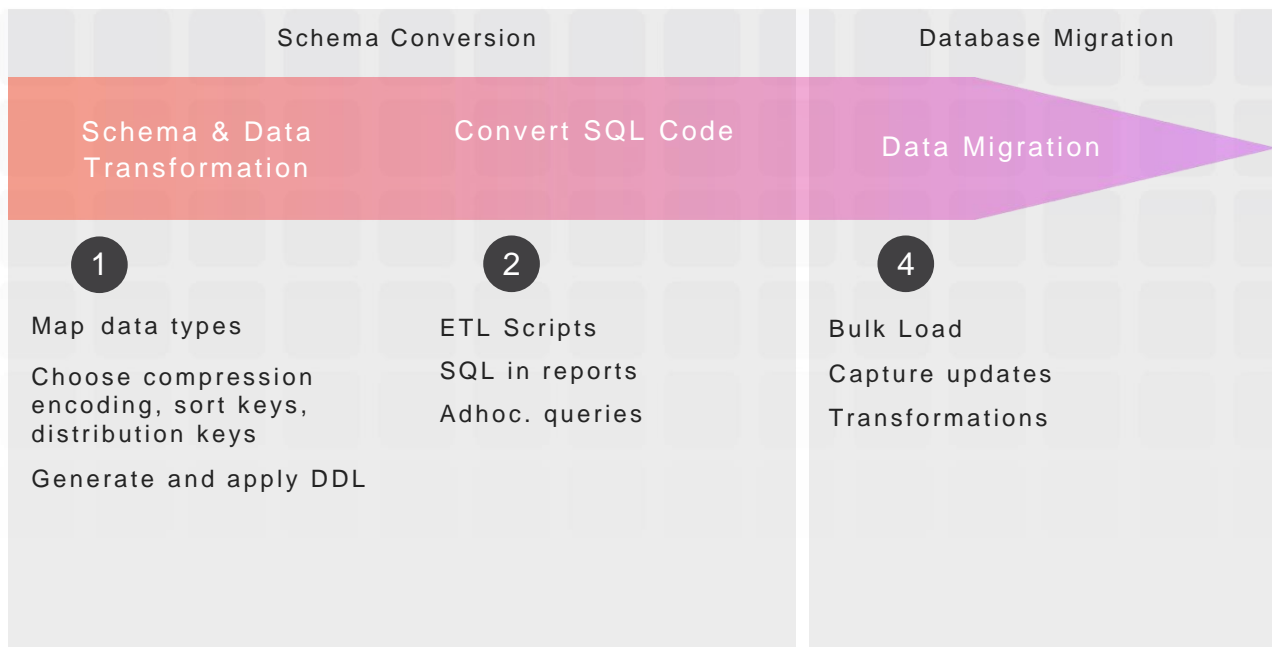
Assess Gaps

Stored Procedures

Functions



ENGINE X



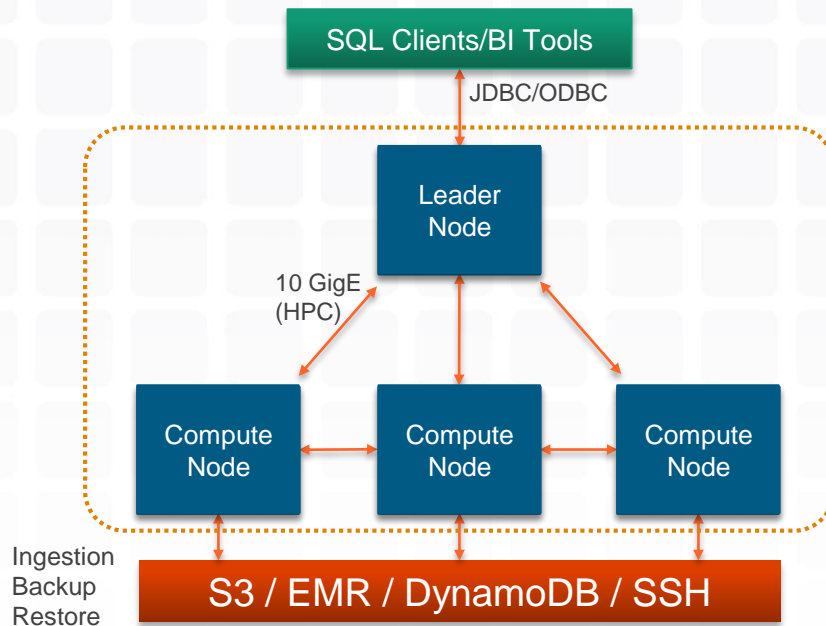
Amazon Redshift

If you forget everything else...

- Lift-and-Shift is **NOT** an ideal approach
 - Depending where you are coming from, it is sure to fail
- AWS has a rich ecosystem of solutions
 - Your final solution will use other AWS services
 - AWS Solution Architects, ProServ, and Partners can help

Amazon Redshift Cluster Architecture

- **Massively parallel, shared nothing**
- **Leader node**
 - SQL endpoint
 - Stores metadata
 - Coordinates parallel SQL processing
- **Compute nodes**
 - Local, columnar storage
 - Executes queries in parallel
 - Load, backup, restore



Designed for I/O Reduction

- Columnar storage
- Data compression
- Zone maps



```
CREATE TABLE loft_migration (  
  aid      INT      --audience_id  
  ,loc     CHAR(3)  --location  
  ,dt      DATE     --date  
);
```

aid	loc	dt
1	SFO	2016-09-01
2	JFK	2016-09-14
3	SFO	2017-04-01
4	JFK	2017-05-14

- Accessing dt with row storage:
 - Need to read everything
 - Unnecessary I/O

Designed for I/O Reduction

- Columnar storage
- Data compression
- Zone maps



```
CREATE TABLE loft_migration (  
  aid      INT      --audience_id  
  ,loc     CHAR(3)  --location  
  ,dt      DATE     --date  
);
```

aid	loc	dt
1	SFO	2016-09-01
2	JFK	2016-09-14
3	SFO	2017-04-01
4	JFK	2017-05-14


- Accessing dt with columnar storage:
 - Only scan blocks for relevant column

Designed for I/O Reduction

- Columnar storage

- Data compression

- Zone maps



```
CREATE TABLE loft_migration (  
  aid      INT      ENCODE LZO  
  ,loc     CHAR(3)  ENCODE BYTEDICT  
  ,dt      DATE     ENCODE RUNLENGTH  
);
```

aid	loc	dt
1	SFO	2016-09-01
2	JFK	2016-09-14
3	SFO	2017-04-01
4	JFK	2017-05-14

- Columns grow and shrink independently
- Effective compression ratios due to like data
- Reduces storage requirements
- Reduces I/O

Designed for I/O Reduction

- Columnar storage
- Data compression
- Zone maps



```
CREATE TABLE loft_migration (  
  aid      INT      --audience_id  
  ,loc     CHAR(3)   --location  
  ,dt      DATE      --date  
);
```

aid	loc	dt
1	SFO	2016-09-01
2	JFK	2016-09-14
3	SFO	2017-04-01
4	JFK	2017-05-14

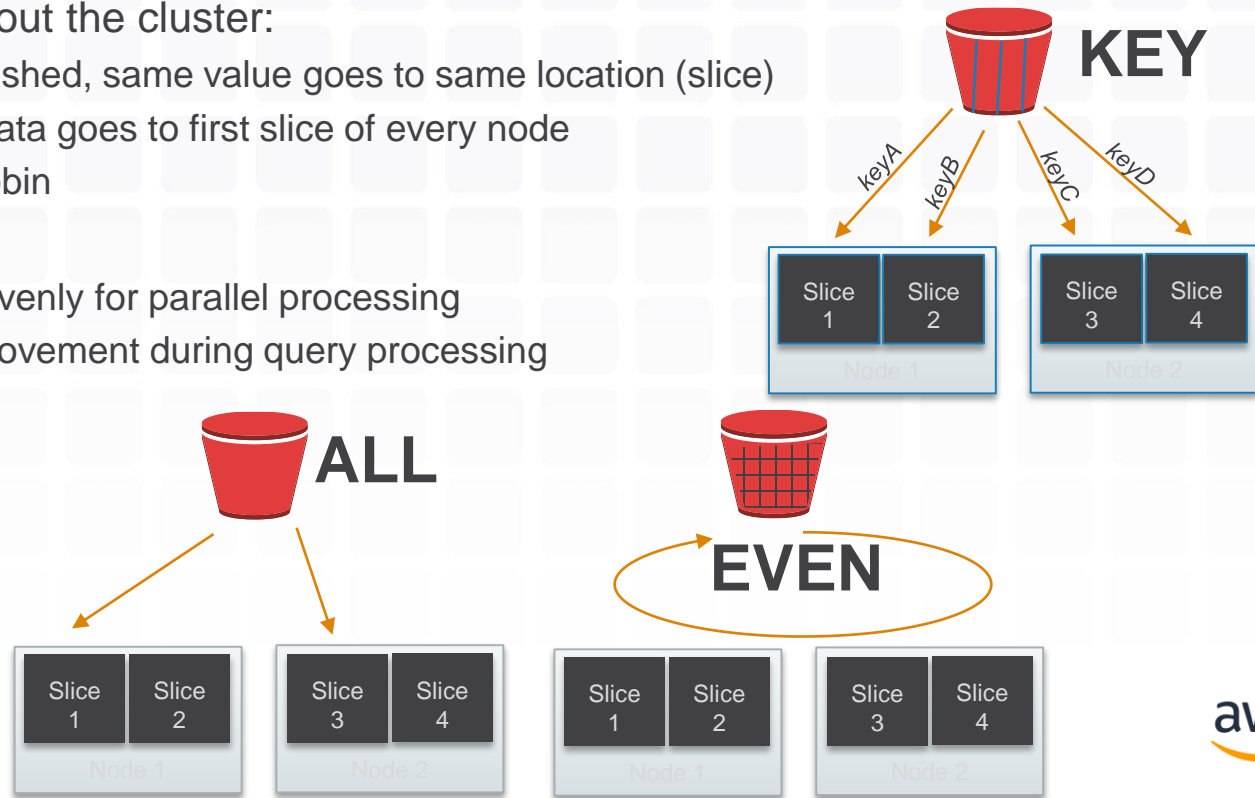
- In-memory block metadata
- Contains per-block MIN and MAX value
- Effectively prunes blocks which cannot contain data for a given query
- Eliminates unnecessary I/O

Terminology and Concepts: Slices

- A **slice** can be thought of like a “virtual compute node”
 - Unit of data partitioning
 - Parallel query processing
- Facts about slices:
 - Each compute node has either 2, 16, or 32 slices
 - Table rows are distributed to slices
 - A slice processes only its own data

Terminology and Concepts: Data Distribution

- **Distribution style** is a table property which dictates how that table's data is distributed throughout the cluster:
 - **KEY:** Value is hashed, same value goes to same location (slice)
 - **ALL:** Full table data goes to first slice of every node
 - **EVEN:** Round robin
- **Goals:**
 - Distribute data evenly for parallel processing
 - Minimize data movement during query processing



Data Loading Best Practices

- Ingestion Throughput:

- Each slice's query processors can load one file at a time:

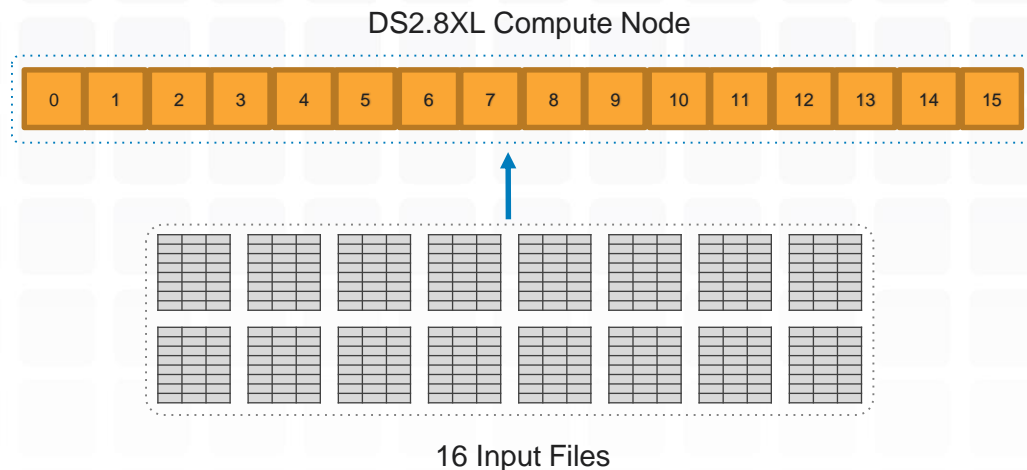
- Streaming decompression
- Parse
- Distribute
- Write



- Realizing only partial node usage as 6.25% of slices are active

Data Loading Best Practices Continued

- Use at least as many input files as there are slices in the cluster
- With 16 input files, all slices are working so you maximize throughput
- COPY continues to scale linearly as you add nodes

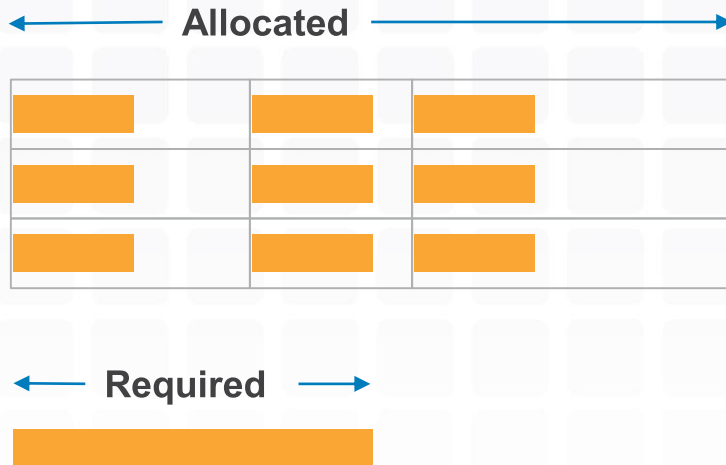


Data Preparation

- Export Data from Source System
 - CSV Recommend (**Delimiter '|'**)
 - Be aware of UTF-8 varchar columns (UTF-8 take 4 bytes per char)
 - Be aware of your NULL character (**\N**)
 - GZIP Compress Files
 - Split Files (1MB – 1GB after gzip compression)
- Useful COPY Options for PoC Data
 - MAXERRORS
 - ACCEPTINVCHARS
 - NULL AS

Keep Columns as Narrow as Possible

- Buffers allocated based on declared column width
- Wider than needed columns mean memory is wasted
- Fewer rows fit into memory; increased likelihood of queries spilling to disk
- Check
`SVV_TABLE_INFO(max_varchar)`
- `SELECT max(len(col)) FROM table`



Amazon Redshift is a Data Warehouse

- Optimized for batch inserts
 - The time to insert a single row in Redshift is roughly the same as inserting 100,000 rows
- Updates are delete + insert of the row
 - Deletes mark rows for deletion
- Blocks are immutable
 - Minimum space used is one block per column, per slice

Column Compression

- Auto Compression
 - Samples data automatically when COPY into an empty table
 - Samples up to 100,000 rows and picks optimal encoding
 - Turn off Auto Compression for Staging Tables
 - Bake encodings into your DDL or use CREATE TABLE (LIKE ...)
- Analyze Compression
 - Data profile has changed
 - Run after changing sort key

Column (Compression) Encoding Types

Raw encoding (RAW)

Byte-dictionary (BYTEDICT)

Delta encoding (DELTA / DELTA32K)

Mostly encoding (MOSTLY8 / MOSTLY16 / MOSTLY32)

Runlength encoding (RUNLENGTH)

Text encoding (TEXT255 / TEXT32K)

LZO encoding (LZO)

Zstandard (ZSTD)

Average: 2-4x

Primary/Unique/Foreign Key Constraints

- Primary/Unique/Foreign Key constraints are NOT enforced
 - If you load data multiple times, Amazon Redshift won't complain
 - If you declare primary keys in your DDL, the optimizer will expect the data to be unique
- Redshift optimizer uses declared constraints to pick optimal plan
 - In certain cases it can result in performance improvements

Benchmarking Tips

- Verify tables are vacuumed and analyzed
 - Check SVV_TABLE_INFO (STATS_OFF, UNSORTED)
 - Vacuum & Analyze
 - <https://github.com/awslabs/amazon-redshift-utils/tree/master/src/AnalyzeVacuumUtility>
- Verify column encodings
 - Check PG_TABLE_DEF
 - Re-encode tables
 - <https://github.com/awslabs/amazon-redshift-utils/tree/master/src/ColumnEncodingUtility>
- Verify good distribution keys
 - SVV_TABLE_INFO (SKEW_ROWS)



AWS Schema Conversion Tool (AWS SCT)

Convert schema in a few clicks

Sources include Oracle, Teradata,
Greenplum, Netezza, and MS SQL DW

Automatic schema optimization

Converts application SQL code

Detailed assessment report

Summary
Action Items

[Save to CSV](#)
[Save to PDF](#)



Database Migration Assessment Report

Source Database

Microsoft SQL Server 2014 - 12.0.4422.0 (X64)
Jul 27 2015 16:56:19
Copyright (c) Microsoft Corporation
Express Edition (64-bit) on Windows NT 6.1 <X64> (Build 7601; Service Pack 1) (Hypervisor)

Executive Summary


We completed the analysis of your SQL Server source database and estimate that 94% of the database storage objects can be converted automatically or with minimal changes if you select MySQL as your migration target. Database storage objects include schemas, tables, columns, constraints, indexes, sequences, synonyms, user-defined types and types. Database code objects include functions, procedures, packages, triggers, views, materialized views, events, SQL

Database Objects with Conversion Actions for MySQL

Of the total 179 database storage object(s) in the source database, we were able to identify 169 (94%) database storage object(s) that can be c

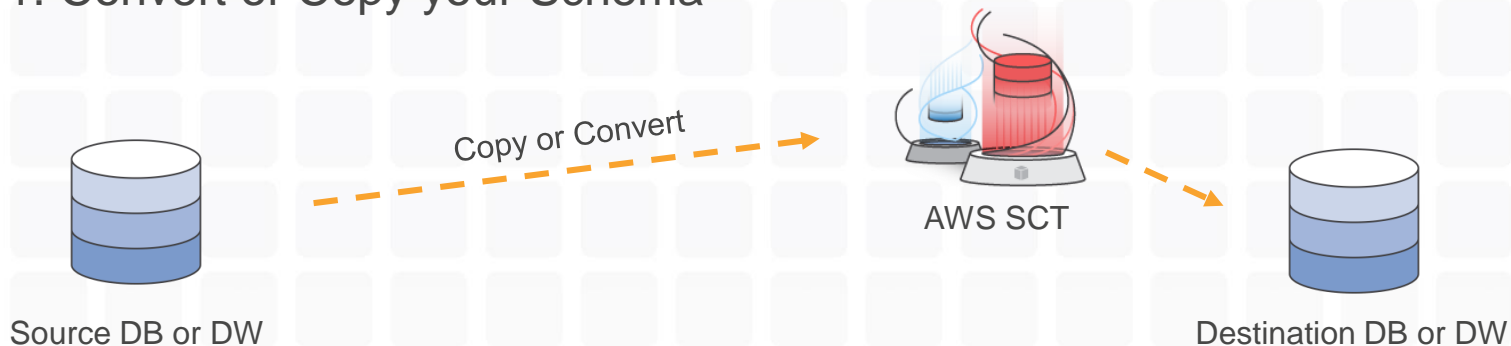
10 (6%) database storage object(s) required 5/8 medium, and 10 significant user action(s) to complete the conversion.

Figure: Conversion statistics for database storage objects

[illegible]

Database migration process

Step 1: Convert or Copy your Schema




Step 2: Move your data



SCT data extractors

Extract Data from your data warehouse and migrate to Amazon Redshift

- Extracts through local migration agents
- Data is optimized for Redshift and Saved in local files
- Files are loaded to an Amazon S3 bucket (through network or Amazon Snowball) and then to Amazon Redshift



For Large Data Transfers

- AWS Snowball
- AWS Snowball Edge
- AWS Snowmobile



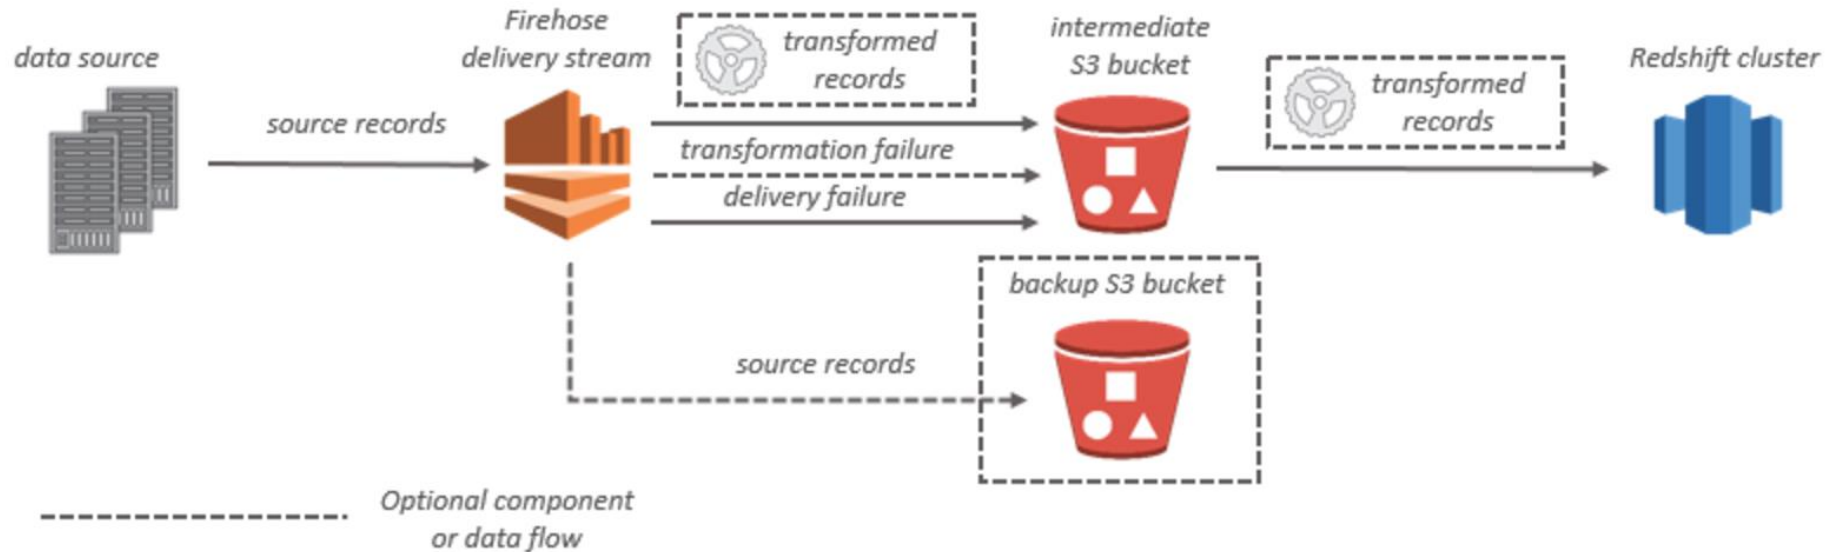
Amazon Kinesis Firehose

Load massive volumes of streaming data into Amazon S3, Redshift and Elasticsearch



- **Zero administration:** Capture and deliver streaming data into Amazon S3, Amazon Redshift, and other destinations **without writing an application or managing infrastructure.**
- **Direct-to-data store integration:** **Batch, compress, and encrypt** streaming data for delivery into data destinations **in as little as 60 secs** using simple configurations.
- **Seamless elasticity:** Seamlessly scales to match data throughput w/o intervention
- **Serverless ETL using AWS Lambda** - Firehose can invoke your Lambda function to transform incoming source data.

Amazon Kinesis Firehose to Amazon Redshift



Data Integration Partners



Amazon Redshift




Data Integration



Systems Integrators



AWS Glue for Automated, Serverless ETL



Data Catalog

Hive metastore compatible metadata repository of data sources
Crawls data source to infer table, data type, partition format



Job Authoring

Generates Python code to move data from source to destination
Edit with your favorite IDE; share code snippets using Git



Job Execution

Runs jobs in Spark containers – automatic scaling based on SLA
Glue is serverless – only pay for the resources you consume

Resources

- <https://github.com/awslabs/amazon-redshift-utils>
- <https://github.com/awslabs/amazon-redshift-monitoring>
- <https://github.com/awslabs/amazon-redshift-udfs>
- **Admin scripts**
Collection of utilities for running diagnostics on your cluster
- **Admin views**
Collection of utilities for managing your cluster, generating schema DDL, etc.
- **ColumnEncodingUtility**
Gives you the ability to apply optimal column encoding to an established schema with data already loaded
- **Amazon Redshift Engineering's Advanced Table Design Playbook**
<https://aws.amazon.com/blogs/big-data/amazon-redshift-engineerings-advanced-table-design-playbook-preamble-prerequisites-and-prioritization/>

Thank you!

briskman@amazon.com

